Linear Regression

Jia-Bin Huang
Virginia Tech

ECE-5424G / CS-5824

Spring 2019

Recap: Machine learning algorithms

	Supervised Learning	**Unsupervised Learning**
Discrete	Classification	Clustering
Continuous	Regression	Dimensionality reduction

Today’s plan: Linear Regression

- Model representation
- Cost function
- Gradient descent
- Features and polynomial regression
- Normal equation

Linear Regression

- **Model representation**
- Cost function
- Gradient descent
- Features and polynomial regression
- Normal equation

<u>Regression</u>

real-valued output

Training set

Learning Algorithm

x → h → y

Size of house

Hypothesis

Estimated price

House pricing prediction

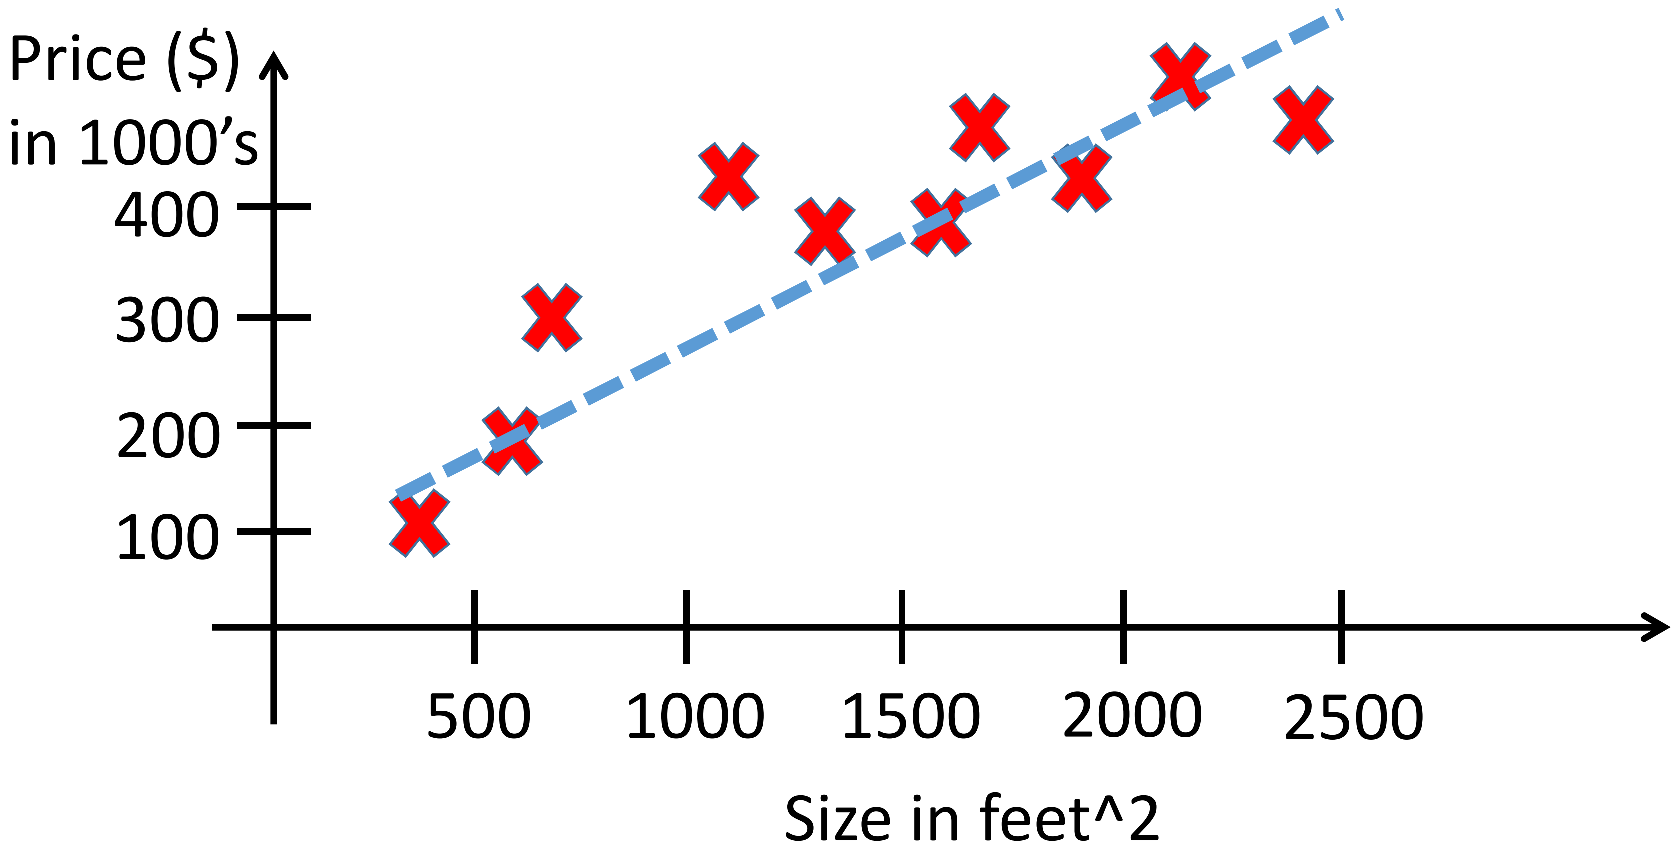

Training set

Size in feet^2 (x)	Price ($) in 1000's (y)
2104	460
1416	232
1534	315
852	178
…	…

m = 47

- Notation:
  - m = Number of training examples
  - x = Input variable / features
  - y = Output variable / target variable
  - (x, y) = One training example
  - $(x^{(i)}, y^{(i)})$ = i^{th} training example

Examples:

$x^{(1)} = 2104$

$x^{(2)} = 1416$

$y^{(1)} = \ 460$

Slide credit: Andrew Ng

Model representation

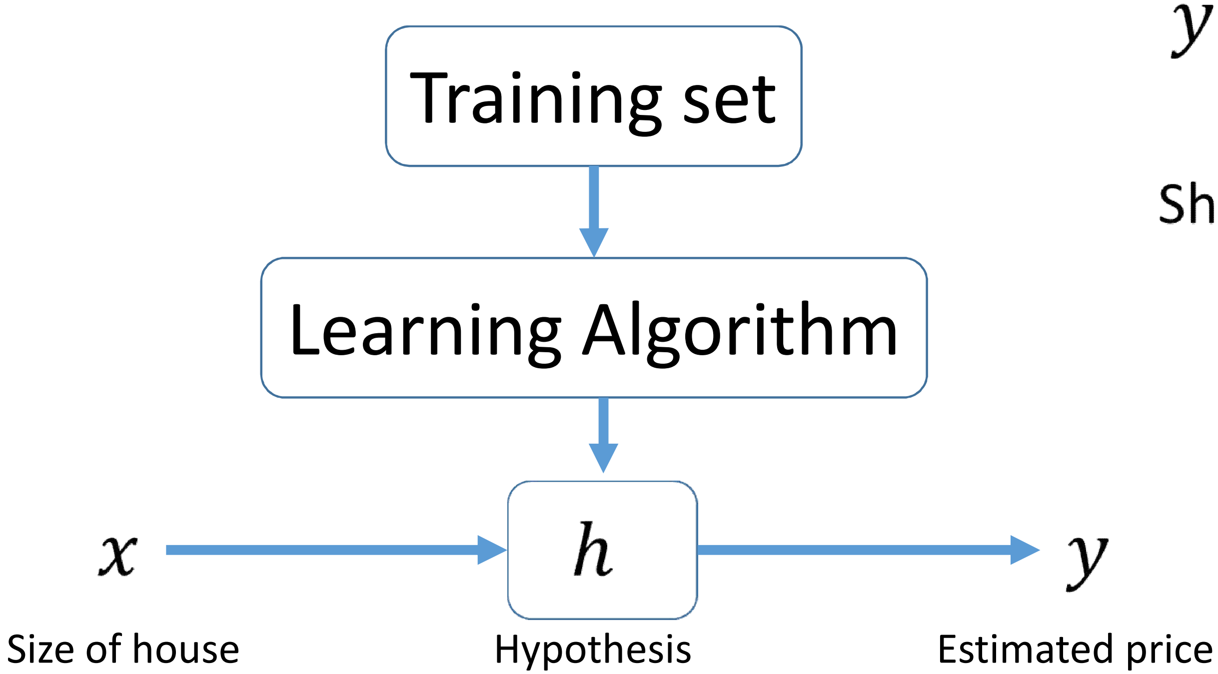

$$y = h_\theta(x) = \theta_0 + \theta_1 x$$

Shorthand $h(x)$

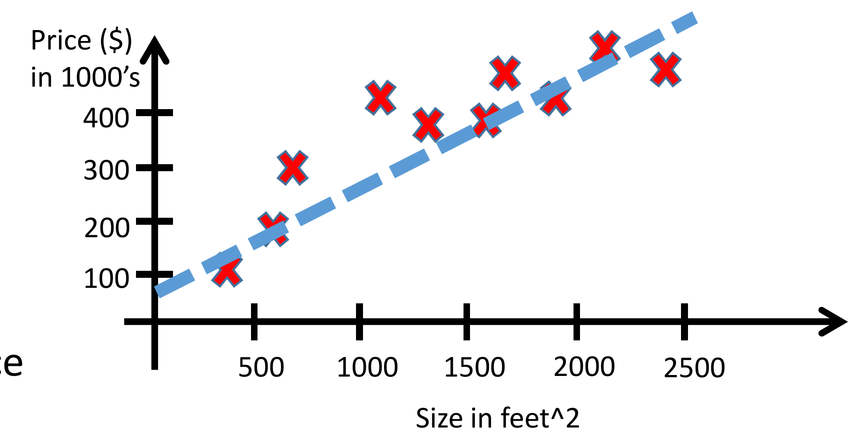

Univariate linear regression

Slide credit: Andrew Ng

Linear Regression

- Model representation
- **Cost function**
- Gradient descent
- Features and polynomial regression
- Normal equation

Training set

Size in feet^2 (x)	Price ($) in 1000's (y)
2104	460
1416	232
1534	315
852	178
…	…

m = 47

- Hypothesis $h_\theta(x) = \theta_0 + \theta_1 x$

θ_0, θ_1: parameters/weights

How to choose θ_i's?

Slide credit: Andrew Ng

$$h_\theta(x) = \theta_0 + \theta_1 x$$

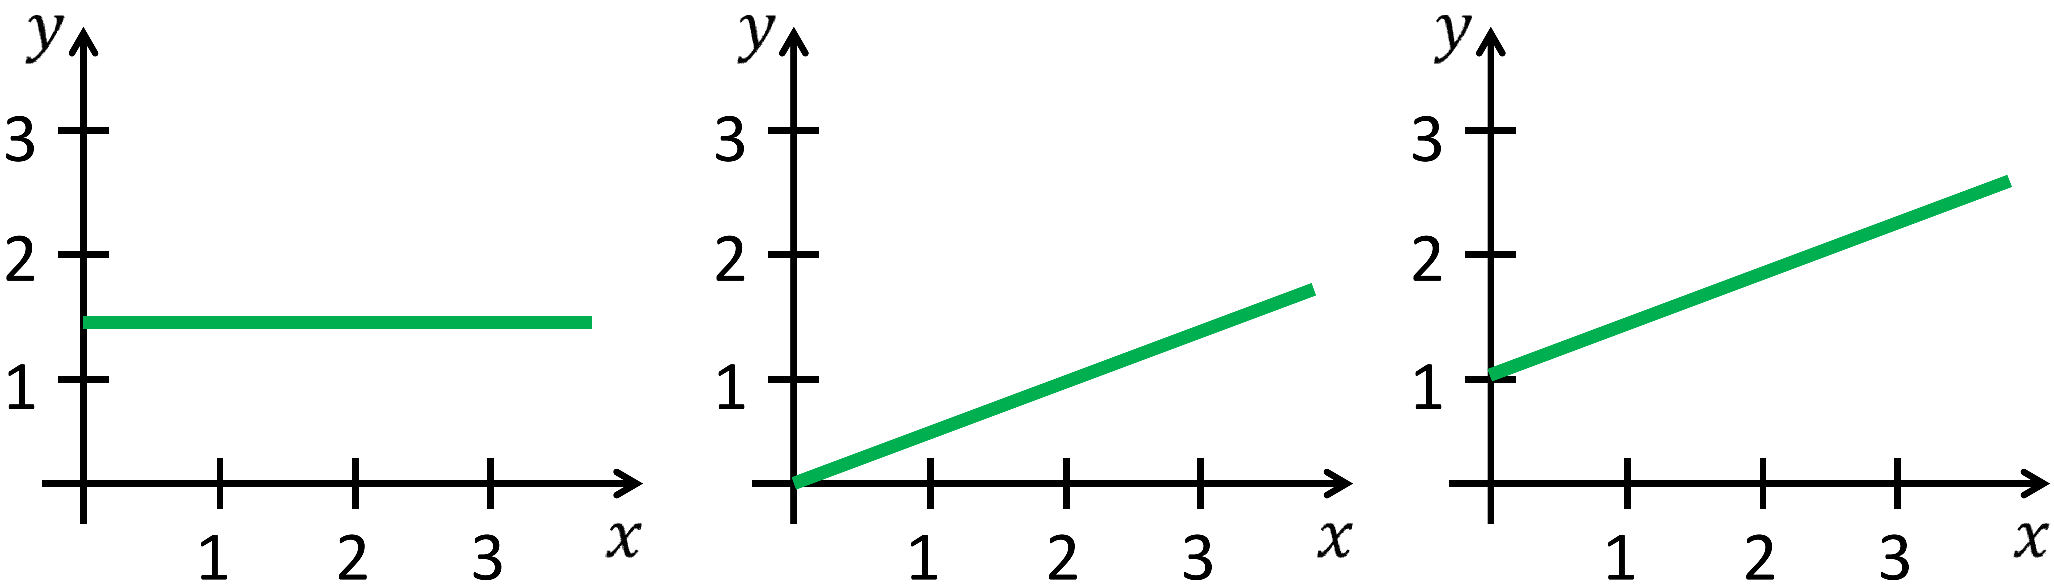

$\theta_0 = 1.5$
$\theta_1 = 0$

$\theta_0 = 0$
$\theta_1 = 0.5$

$\theta_0 = 1$
$\theta_1 = 0.5$

Slide credit: Andrew Ng

Cost function

- Idea:
  Choose θ_0, θ_1 so that $h_\theta(x)$ is close to y for our training example (x, y)

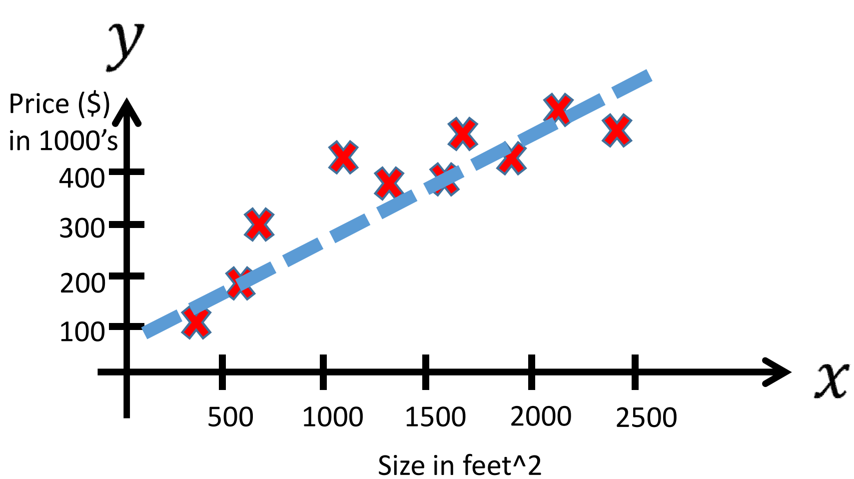

$$\underset{\theta_0, \theta_1}{\text{minimize}} \ \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$$

$$h_\theta\left(x^{(i)}\right) = \theta_0 + \theta_1 x^{(i)}$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$$

$$\underset{\theta_0, \theta_1}{\text{minimize}} \ \boxed{J(\theta_0, \theta_1)}$$ **Cost function**

Slide credit: Andrew Ng

Simplified

- **Hypothesis:**

$$h_\theta(x) = \theta_0 + \theta_1 x \longrightarrow h_\theta(x) = \theta_1 x \qquad \theta_0 = 0$$

- **Parameters:**

$$\theta_0, \theta_1 \longrightarrow \theta_1$$

- **Cost function:**

$$J(\theta_0, \theta_1) = \frac{1}{2m}\sum_{i=1}^{m}\left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2 \longrightarrow J(\theta_1) = \frac{1}{2m}\sum_{i=1}^{m}\left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$$

- **Goal**:

$$\underset{\theta_0, \theta_1}{\text{minimize}}\ J(\theta_0, \theta_1) \longrightarrow \underset{\theta_0, \theta_1}{\text{minimize}}\ J(\theta_1)$$

Slide credit: Andrew Ng

$h_\theta(x)$, function of x

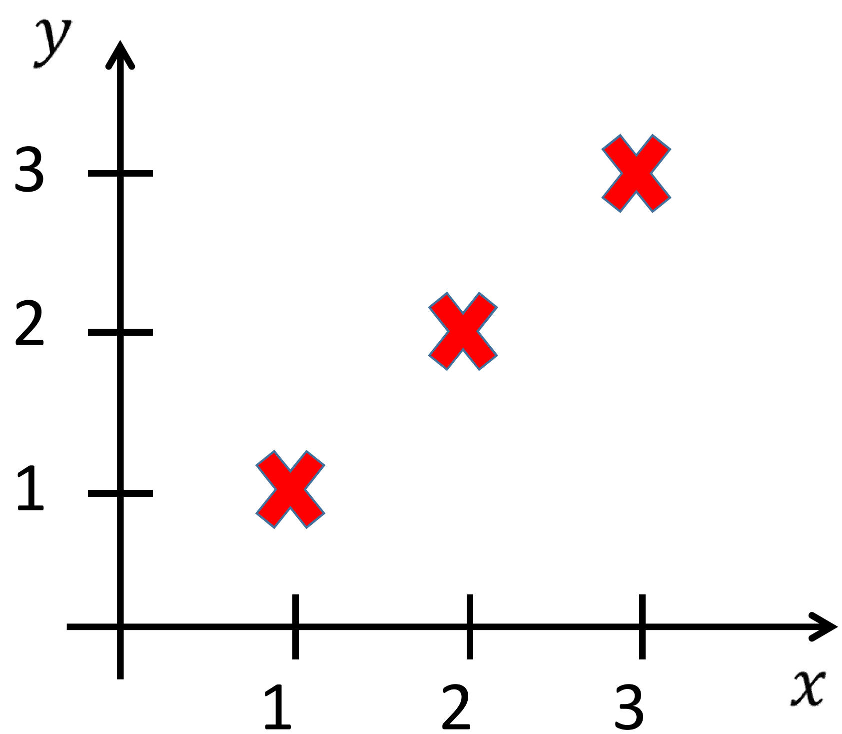

$J(\theta_1)$, function of θ_1

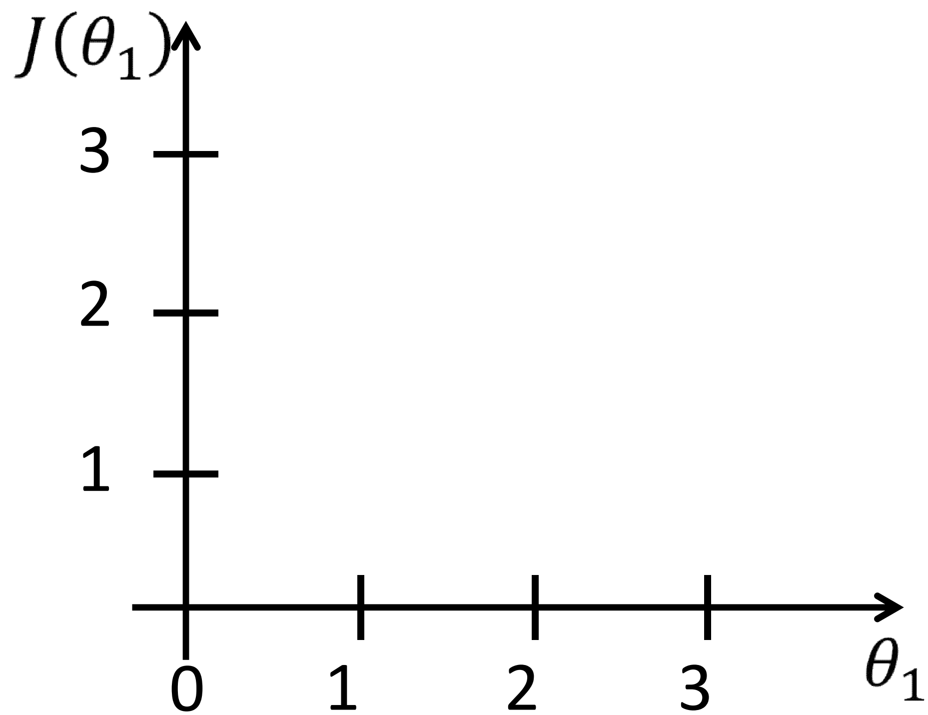

Slide credit: Andrew Ng

$h_\theta(x)$, function of x

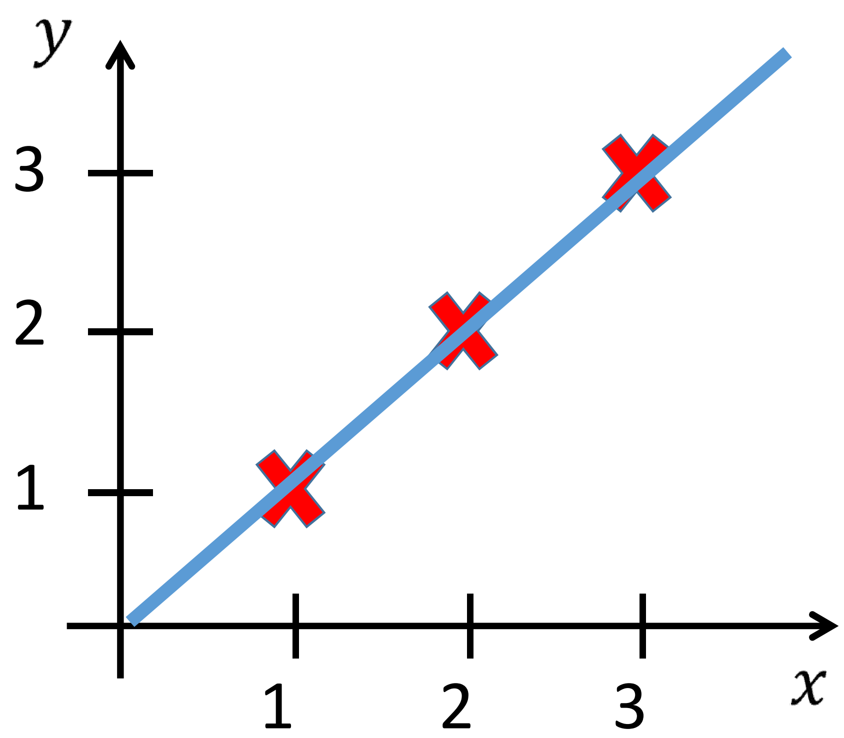

$J(\theta_1)$, function of θ_1

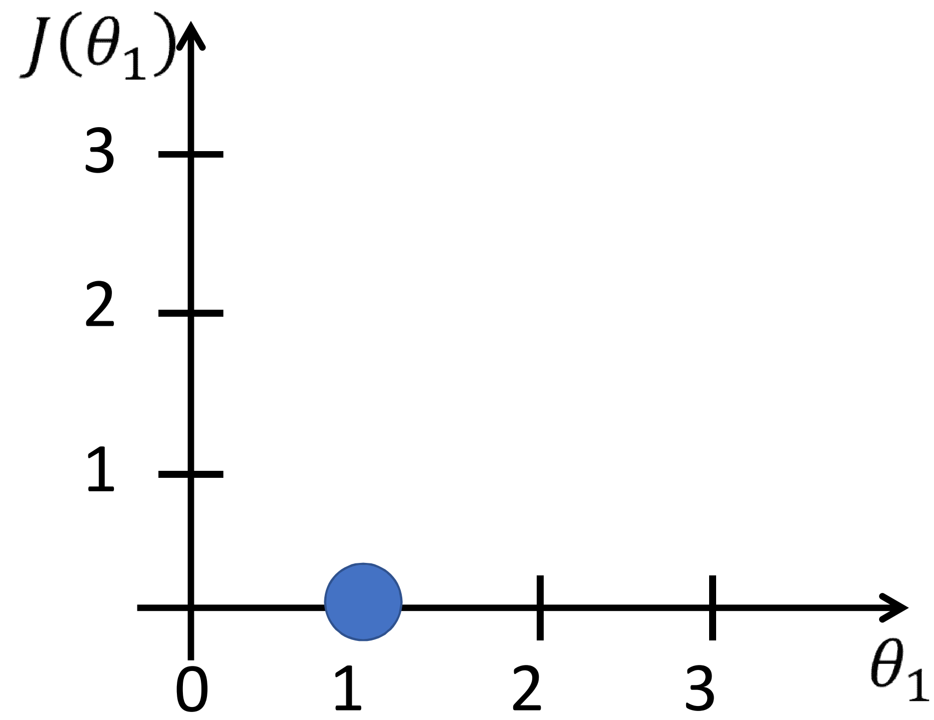

Slide credit: Andrew Ng

$h_\theta(x)$, function of x

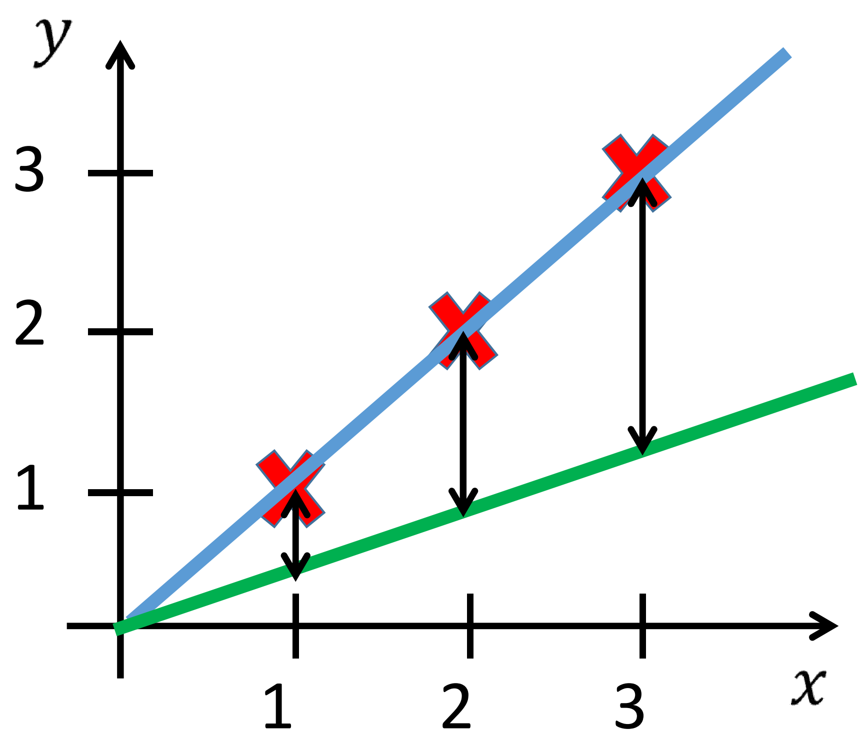

$J(\theta_1)$, function of θ_1

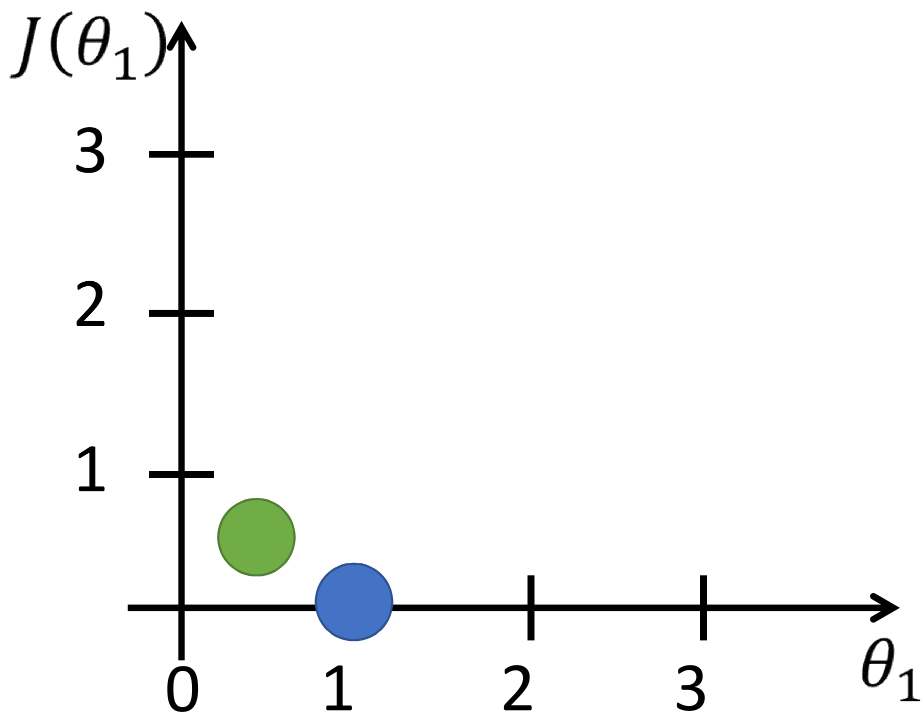

Slide credit: Andrew Ng

$h_\theta(x)$, function of x

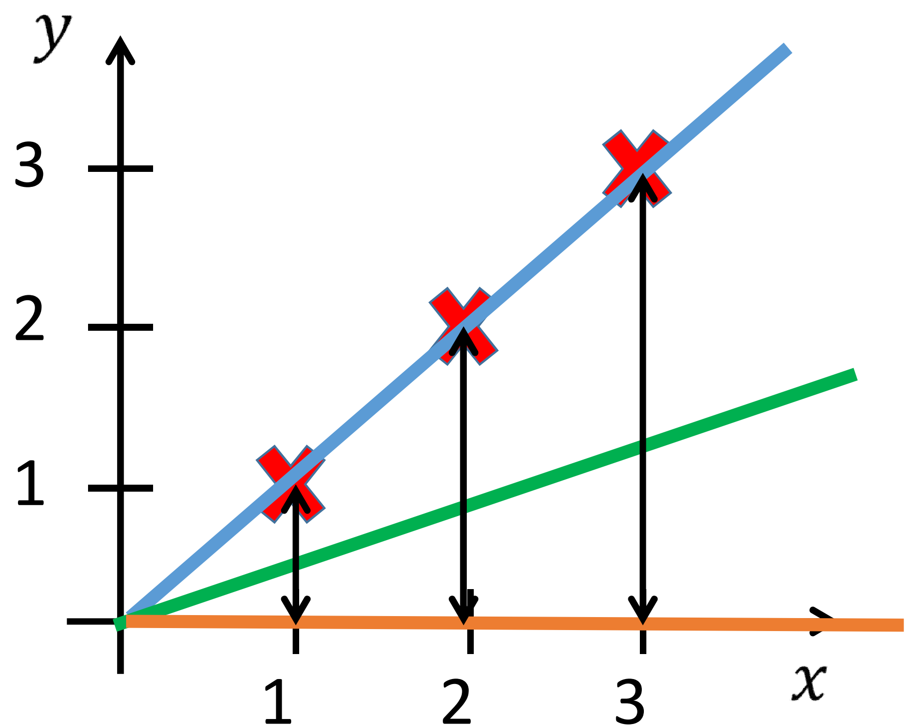

$J(\theta_1)$, function of θ_1

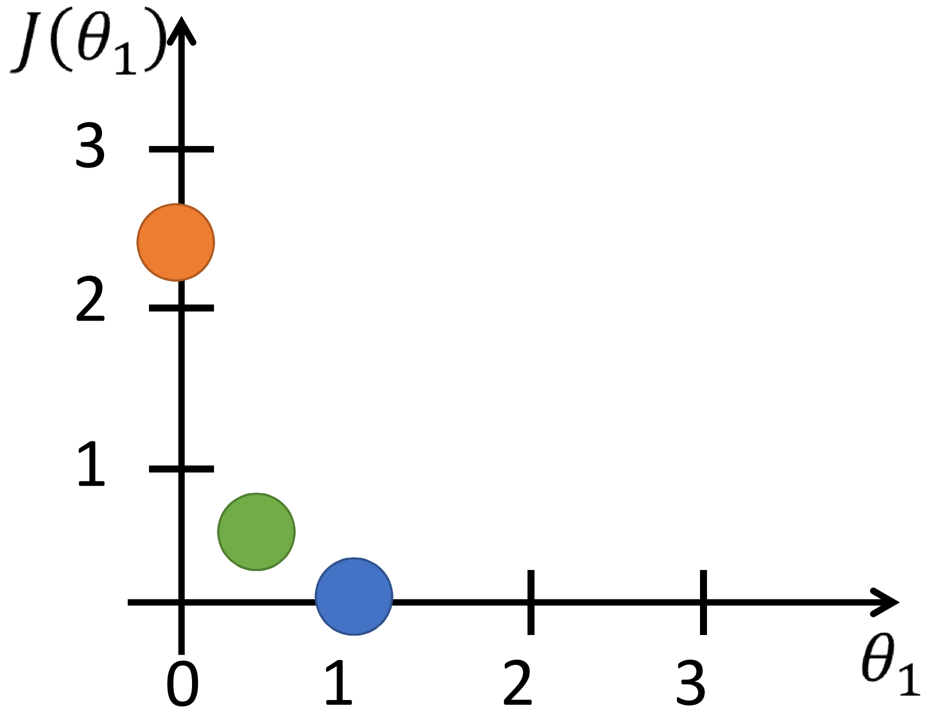

Slide credit: Andrew Ng

$h_\theta(x)$, function of x

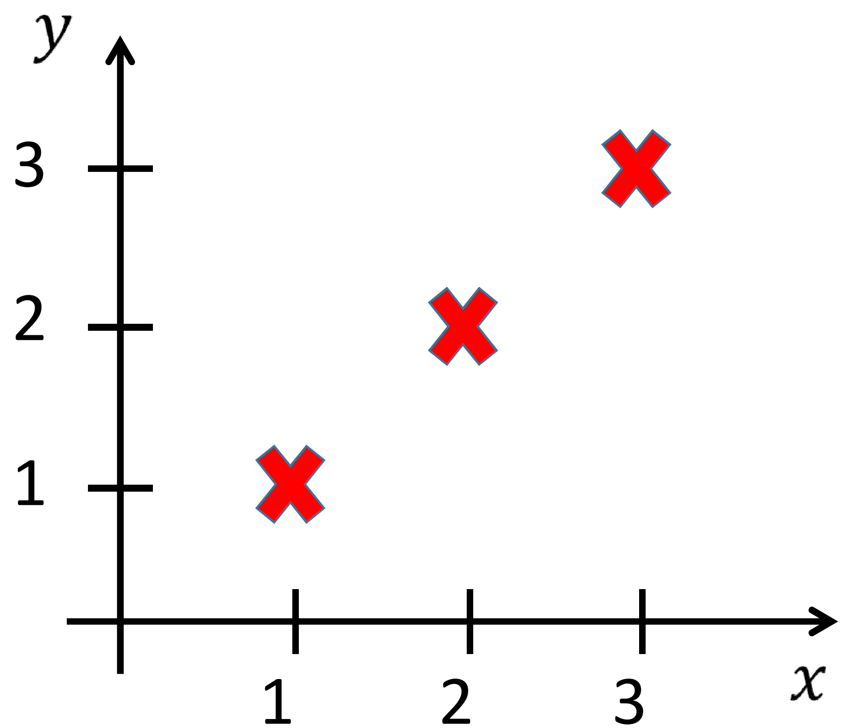

$J(\theta_1)$, function of θ_1

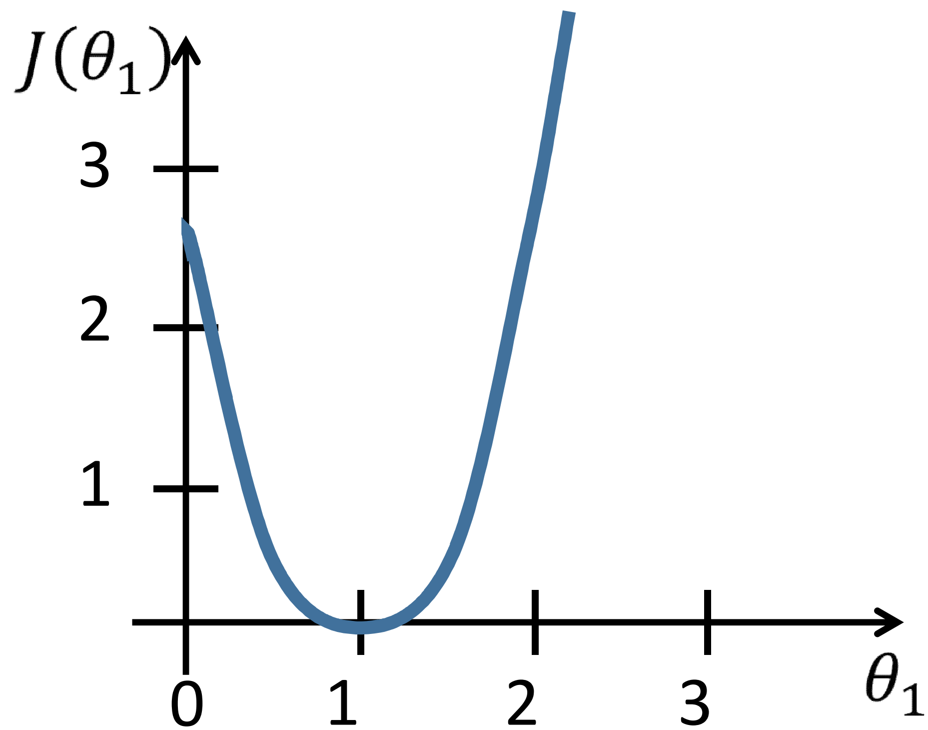

Slide credit: Andrew Ng

- **Hypothesis:** $h_\theta(x) = \theta_0 + \theta_1 x$

- **Parameters:** θ_0, θ_1

- **Cost function:** $J(\theta_0, \theta_1) = \frac{1}{2m}\sum_{i=1}^{m}\left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$

- **Goal**: $\underset{\theta_0, \theta_1}{\text{minimize}}\; J(\theta_0, \theta_1)$

Slide credit: Andrew Ng

Cost function

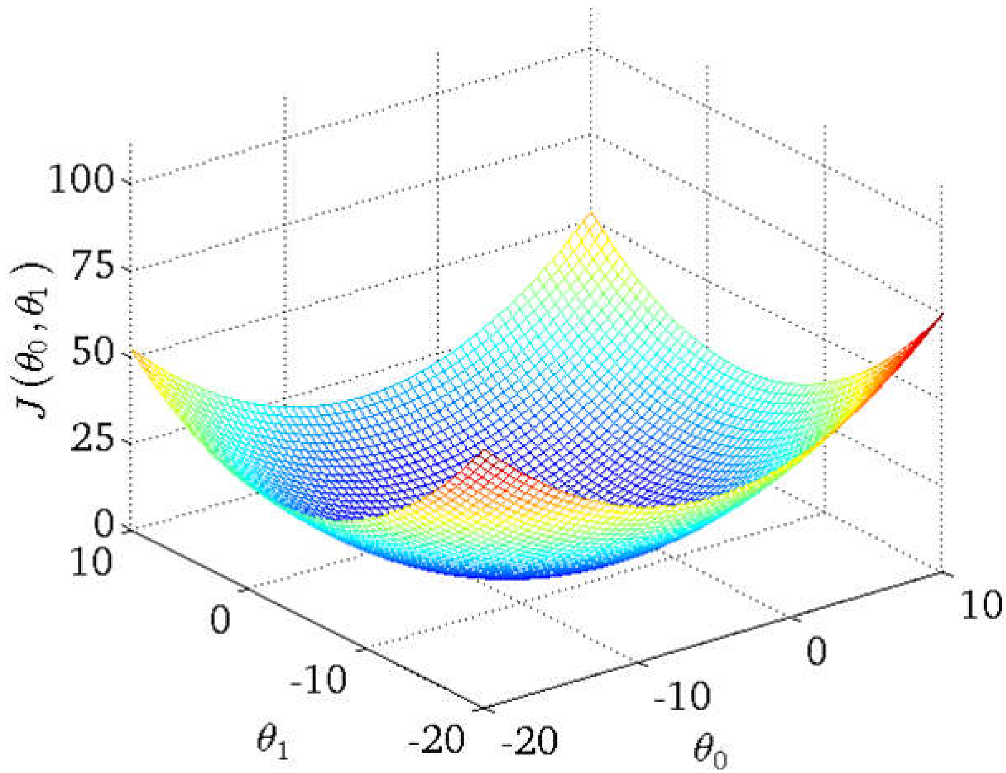

Slide credit: Andrew Ng

What are Contour Plots?

- Contour plots are a way to show a three-dimensional surface on a two-dimensional plane. It graphs two predictor variables X, Y on the y-axis and a response variable Z as contours. These contours are sometimes called *z-slices* or *iso-response values*.
- A contour plot is appropriate if you want to see how some value Z changes as a function of two inputs, X and Y: z = f(x, y).

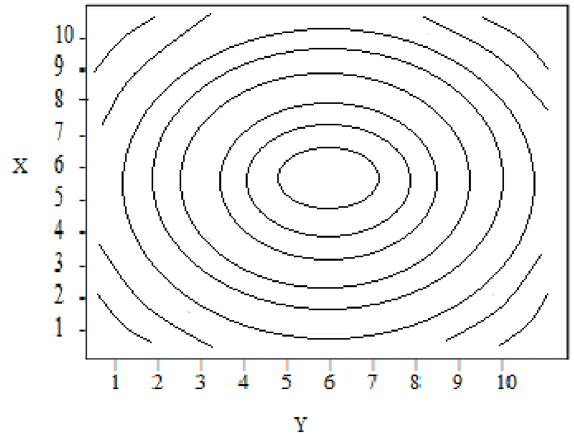

$h_\theta(x)$

(for fixed θ_0, θ_1, this is a function of x)

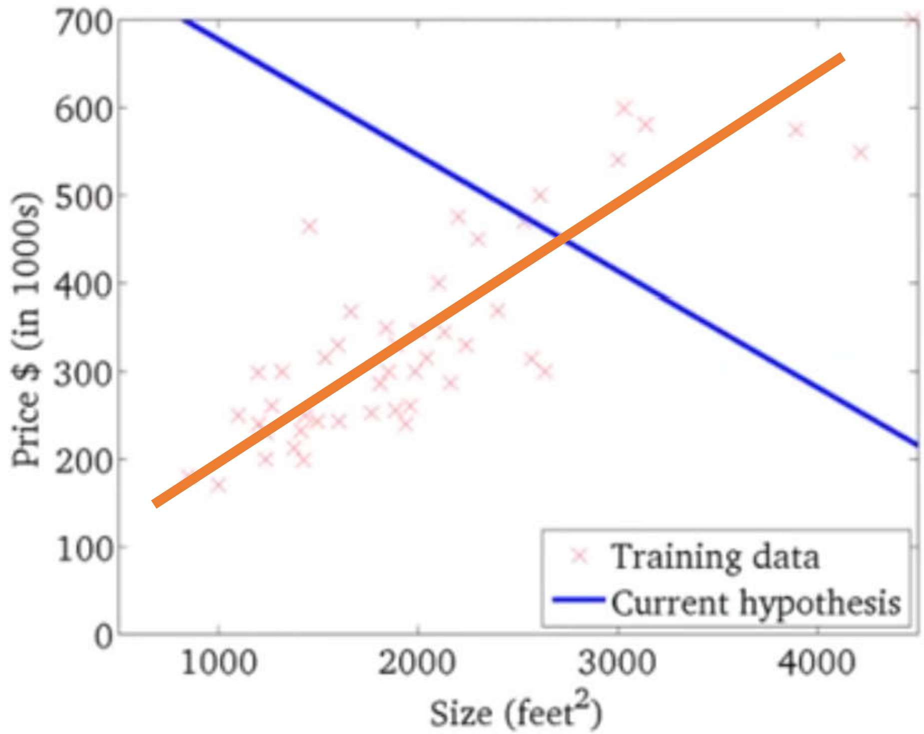

$J(\theta_0, \theta_1)$

(function of the parameters θ_0, θ_1)

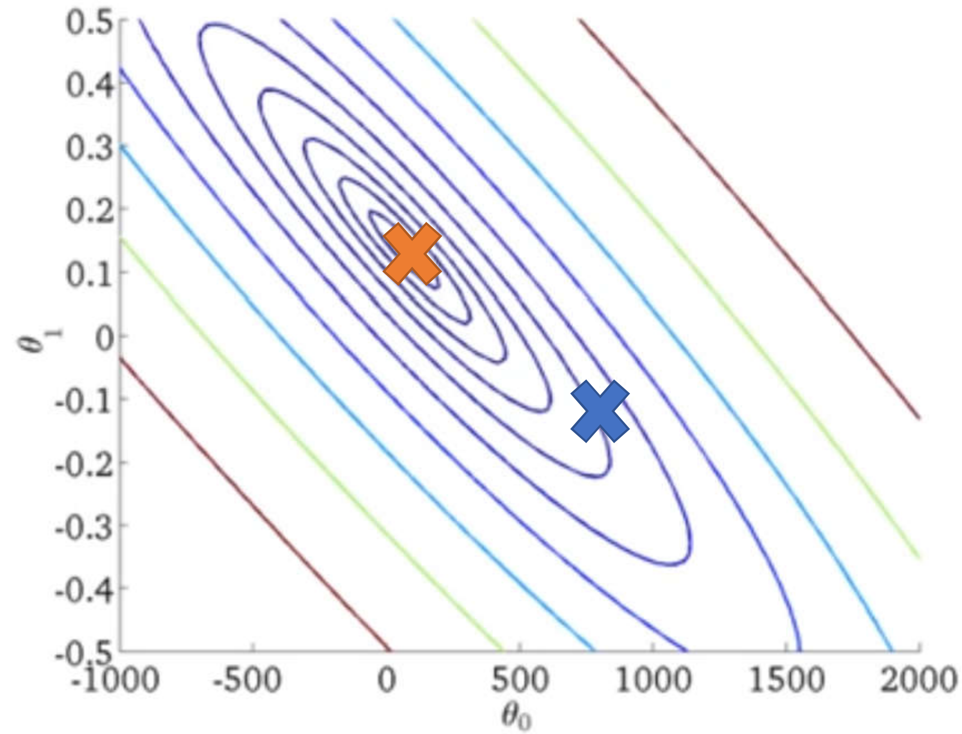

How do we find good θ_0, θ_1 that minimize $J(\theta_0, \theta_1)$?

Slide credit: Andrew Ng

Linear Regression

- Model representation
- Cost function
- **Gradient descent**
- Features and polynomial regression
- Normal equation

Gradient descent

Have some function $J(\theta_0, \theta_1)$

Want $\underset{\theta_0, \theta_1}{\operatorname{argmin}} \; J(\theta_0, \theta_1)$

Outline:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$ until we hopefully end up at minimum

Slide credit: Andrew Ng

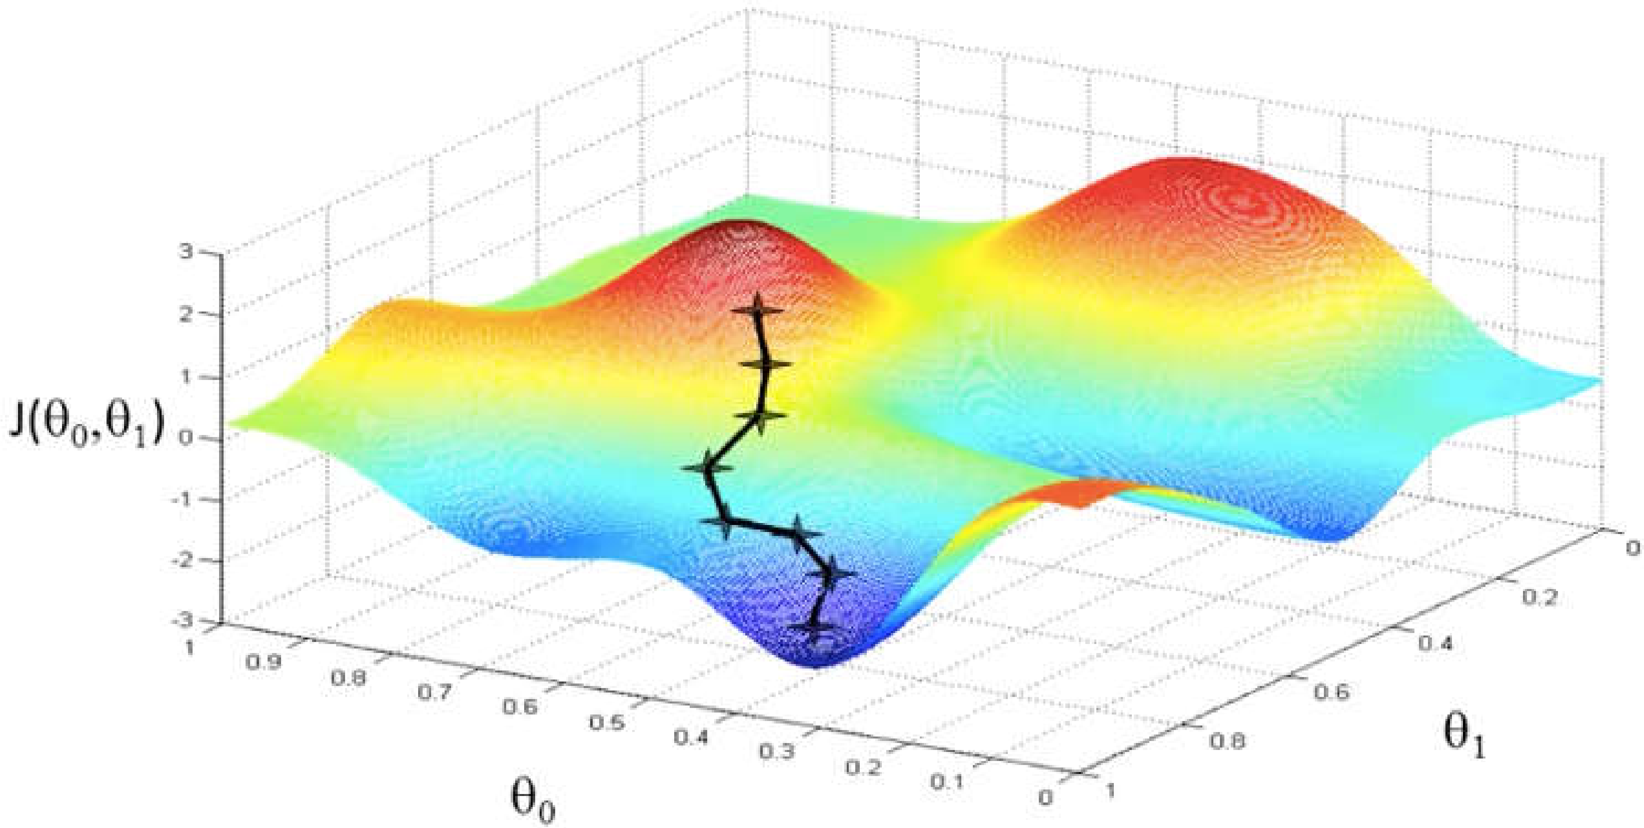

Slide credit: Andrew Ng

Gradient descent

Repeat until convergence{

$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)

}

α: Learning rate (step size)

$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$: derivative (rate of change)

Slide credit: Andrew Ng

Gradient descent

Correct: simultaneous update

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$

Incorrect:

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_1 := \text{temp1}$$

Slide credit: Andrew Ng

$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

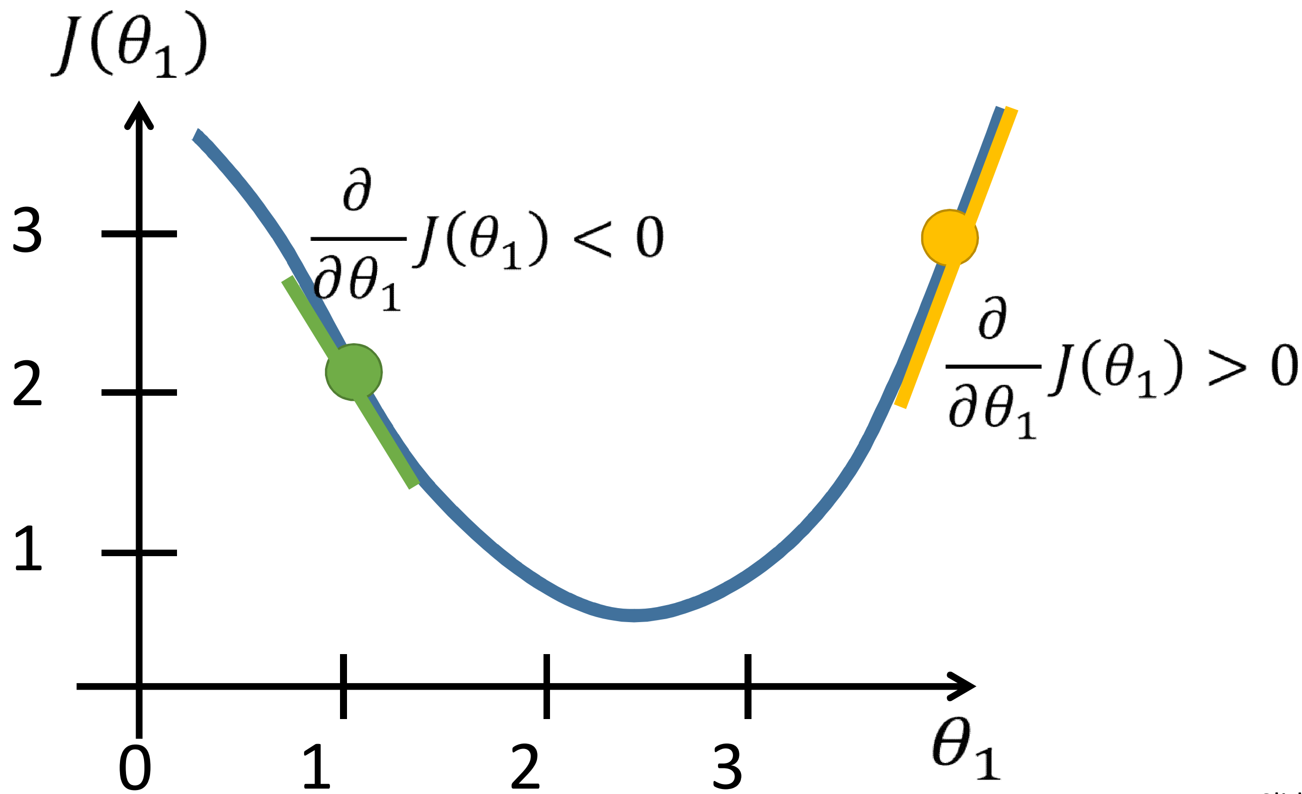

Slide credit: Andrew Ng

Learning rate

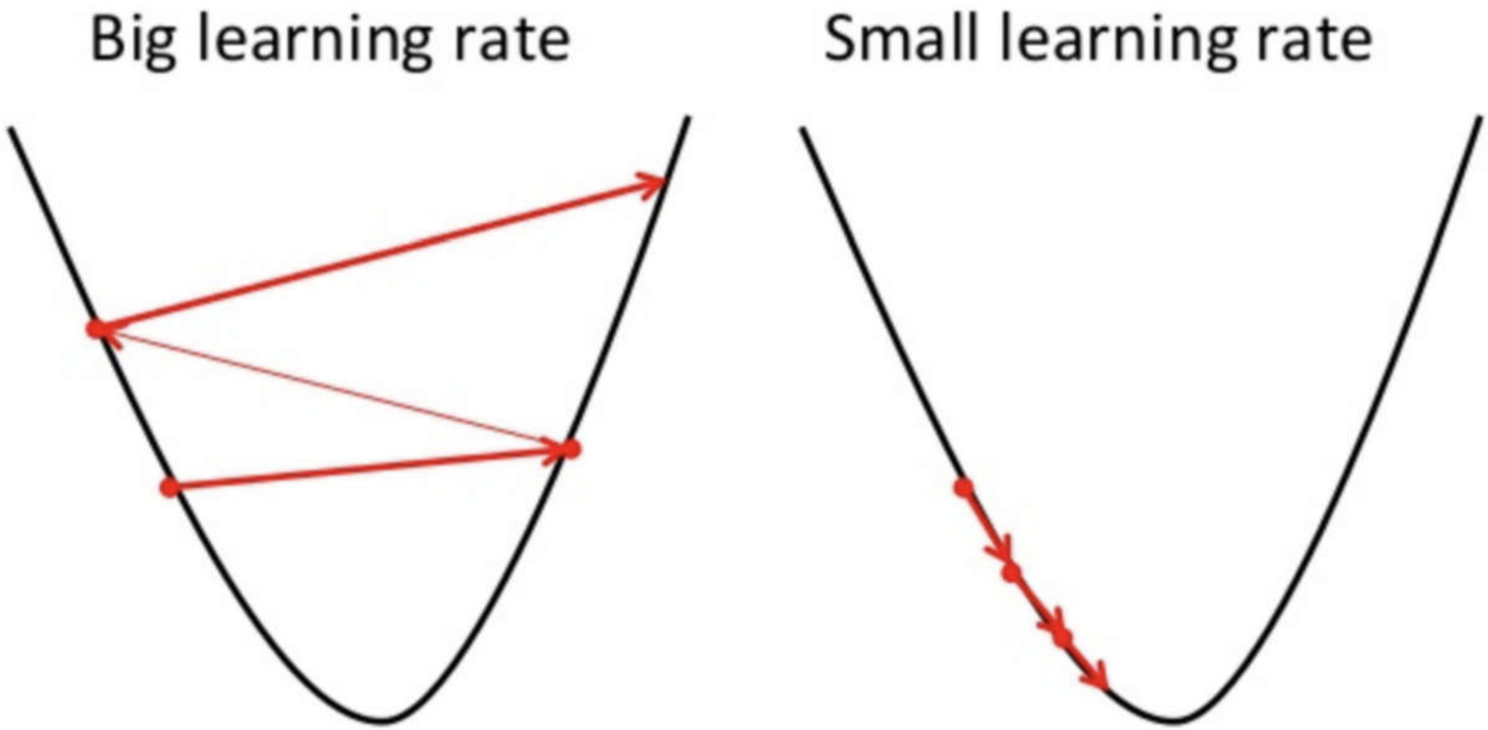

Gradient descent for linear regression

Repeat until convergence{

$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)

}

- Linear regression model

$h_\theta(x) = \theta_0 + \theta_1 x$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$$

Slide credit: Andrew Ng

Computing partial derivative

- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)} \right)^2$

  $= \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^{m} \left(\theta_0 + \theta_1 x^{(i)} - y^{(i)} \right)^2$

- $j = 0$: $\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)} \right)$
- $j = 1$: $\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)} \right) x^{(i)}$

Slide credit: Andrew Ng

Gradient descent for linear regression

Repeat until convergence{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)} \right)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)} \right) x^{(i)}$$

}

Update θ_0 and θ_1 simultaneously

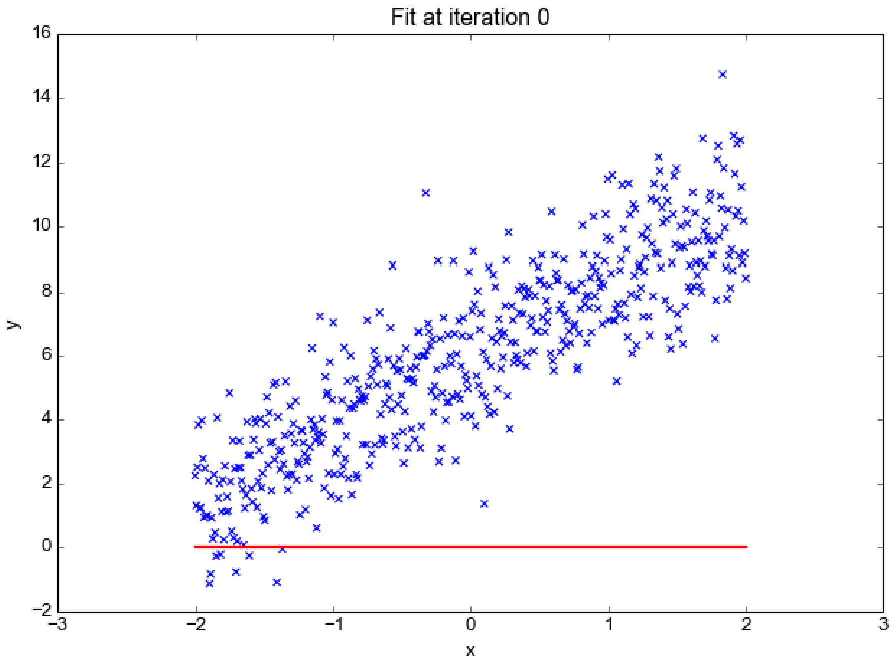

Slide credit: Andrew Ng

Batch gradient descent

- "Batch": Each step of gradient descent uses all the training examples

Repeat until convergence{

m: Number of training examples

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^{\boxed{m}} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^{\boxed{m}} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right) x^{(i)}$$

}

Slide credit: Andrew Ng

Mini Batch gradient descent

Say $b = 10, m = 1000.$

Repeat {

for $i = 1, 11, 21, 31, \ldots, 991$ {

$$\theta_j := \theta_j - \alpha \frac{1}{10} \sum_{k=i}^{i+9} (h_\theta(x^{(k)}) - y^{(k)}) x_j^{(k)}$$

(for every $j = 0, \ldots, n$)

}

}

Mini Batch gradient descent

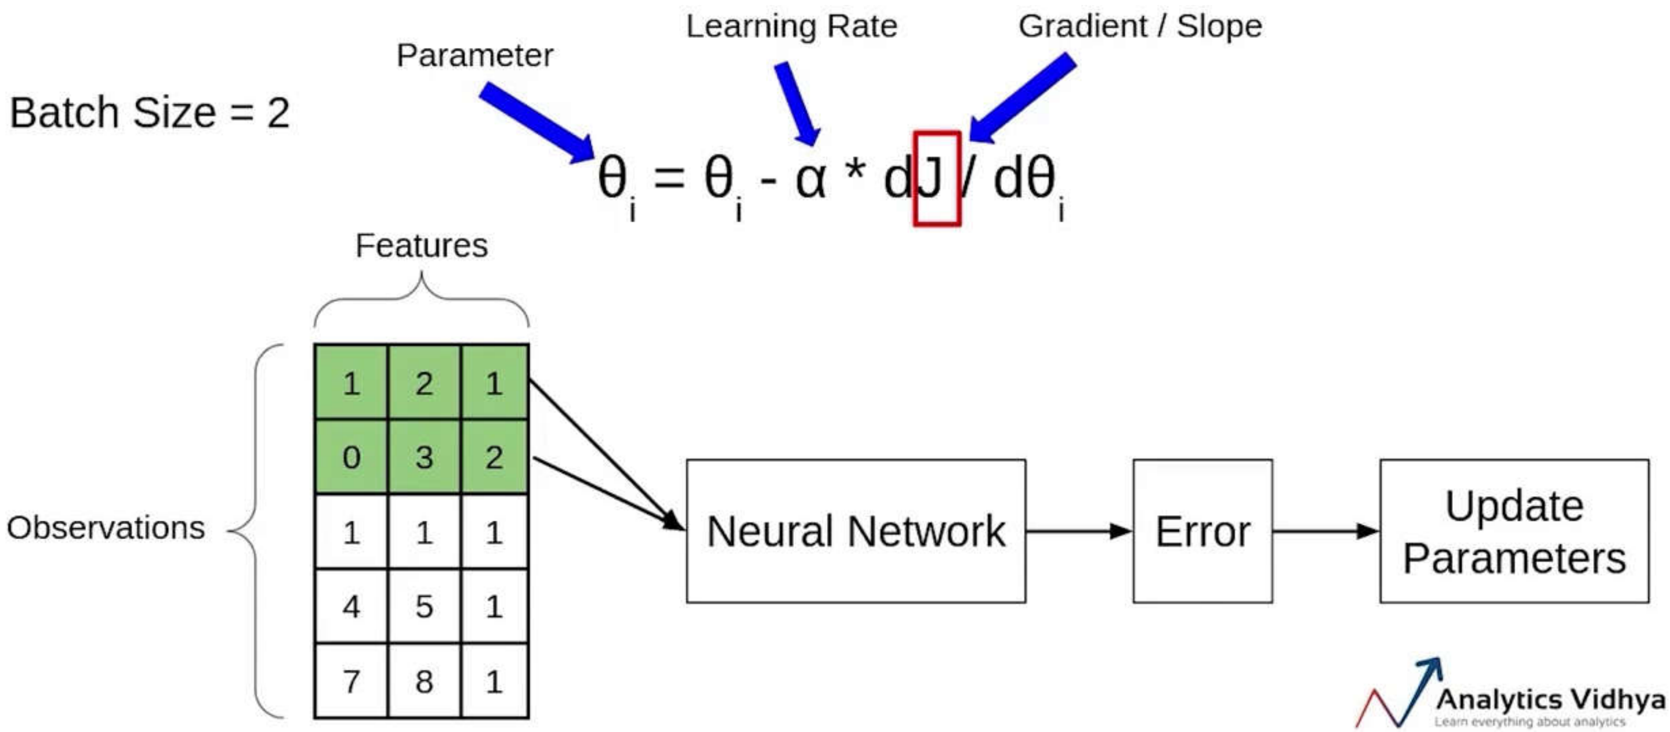

Stochastic Batch gradient descent

Algorithm 2: Pseudo-code for SGD

```
Function SGD:
    Set epsilon as the limit of convergence
    for i = 1, ..., m do
        for j = 0, ..., n do
            while |ω_{j+1} − ω_j| < epsilon do
                ω_{j+1} := ω_j − α · (h_ω(x^(i)) − y^(i)) · x_j^(i);
            end
        end
    end
```

Function SGD:

Set epsilon as the limit of convergence

for $i = 1, \ldots, m$ **do**

for $j = 0, \ldots, n$ **do**

while $|\omega_{j+1} - \omega_j| < epsilon$ **do**

$\omega_{j+1} := \omega_j - \alpha \cdot (h_\omega(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)};$

end

end

end

GD VS MGD VS SGD

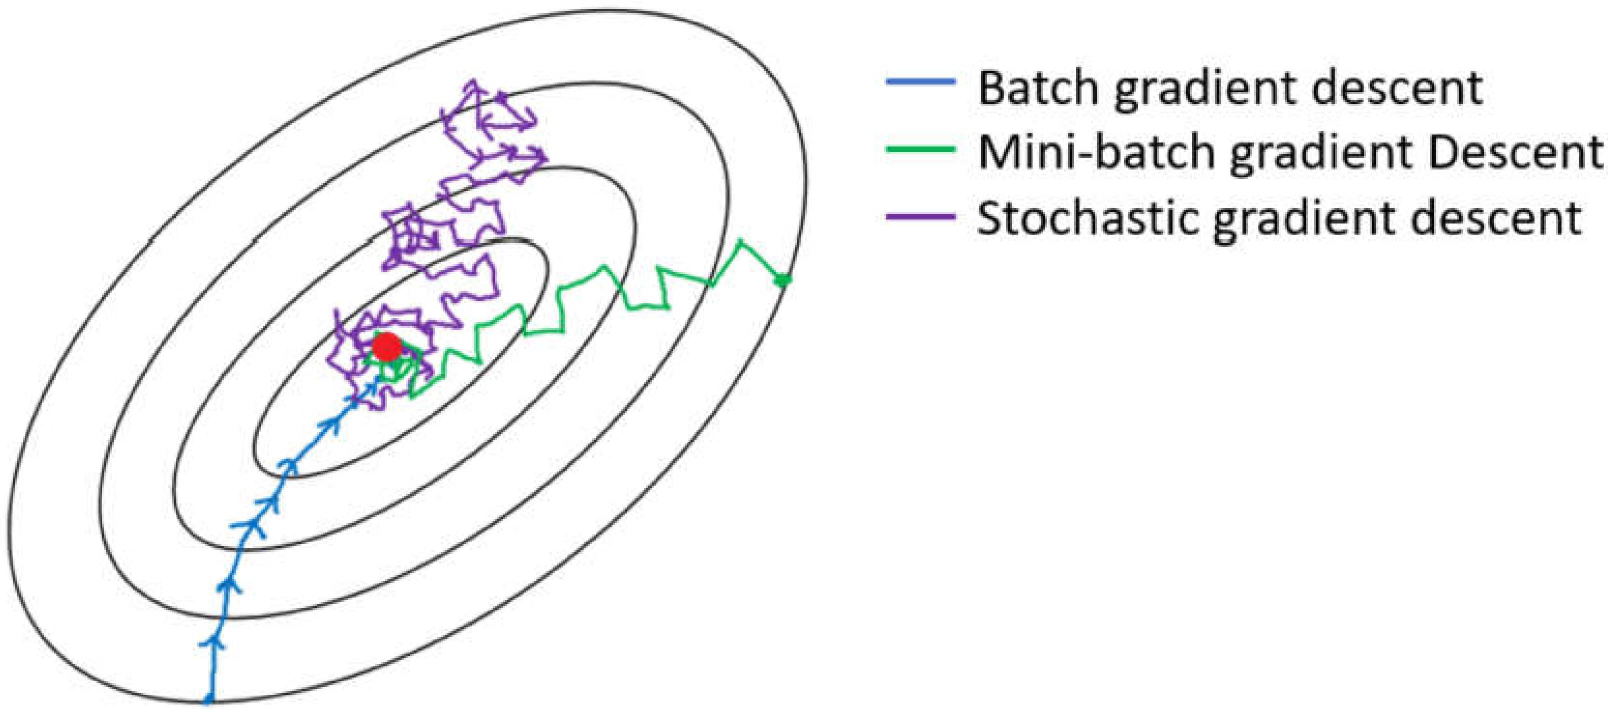

Linear Regression

- Model representation
- Cost function
- Gradient descent
- **Features and polynomial regression**
- Normal equation

Training dataset

Size in feet^2 (x)	Price ($) in 1000's (y)
2104	460
1416	232
1534	315
852	178
…	…

$$h_\theta(x) = \theta_0 + \theta_1 x$$

Slide credit: Andrew Ng

Multiple features (input variables)

Size (feet²) (x_1)	Number of bedrooms (x_2)	Number of floors (x_3)	Age of home (years) (x_4)	Price ($) in 1000's (y)
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
…				…

Notation:

n = Number of features

$x^{(i)}$= Input features of i^{th} training example

$x_j^{(i)}$= Value of feature j in i^{th} training example

$x_3^{(2)} = ?$

$x_3^{(4)} = ?$

Slide credit: Andrew Ng

Hypothesis

Previously: $h_\theta(x) = \theta_0 + \theta_1 x$

Now: $h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$

Slide credit: Andrew Ng

$$h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- For convenience of notation, define $x_0 = 1$ ($x_0^{(i)} = 1$ for all examples)

- $\boldsymbol{x} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in R^{n+1}$ $\quad\quad$ $\boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in R^{n+1}$

- $h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$
  $= \boldsymbol{\theta}^\top \boldsymbol{x}$

Slide credit: Andrew Ng

Gradient descent

- Previously ($n = 1$)

Repeat until convergence{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right) x^{(i)}$$

}

- New algorithm ($n \geq 1$)

Repeat until convergence{

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right) x_j^{(i)}$$

}

Simultaneously update
θ_j, for $j = 0, 1, \cdots, n$

Slide credit: Andrew Ng

Gradient descent in practice: Feature scaling

- Idea: Make sure features are on a similar scale (e.g,. $-1 \leq x_i \leq 1$)
- E.g. x_1 = size (0-2000 feat^2)

  x_2 = number of bedrooms (1-5)

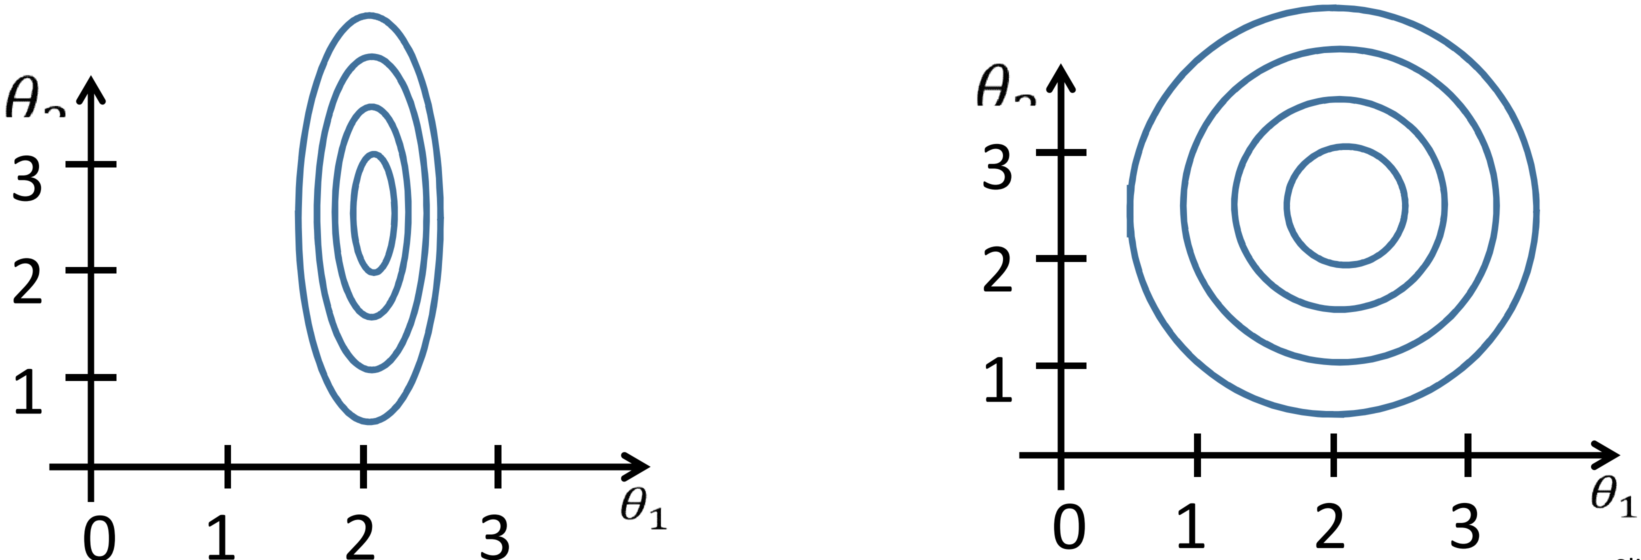

Slide credit: Andrew Ng

Gradient descent in practice: Learning rate

- Automatic convergence test
- α too small: slow convergence
- α too large: may not converge

- To choose α, try

0.001, … 0.01, …, 0.1, … , 1

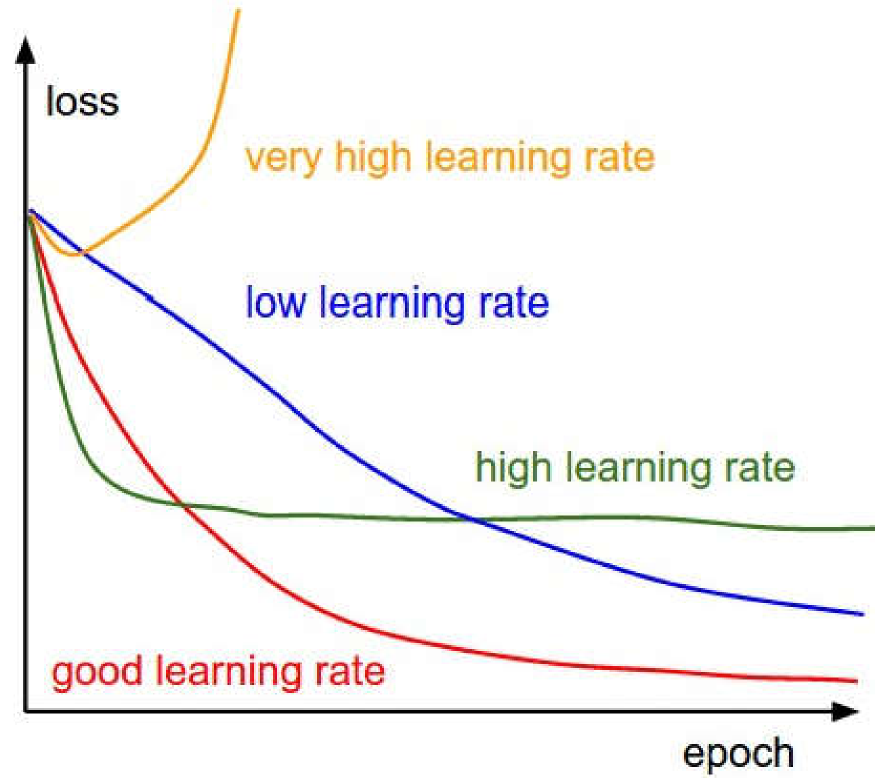

Image credit: CS231n@Stanford

House prices prediction

- $h_\theta(x) = \theta_0 + \theta_1 \times \text{frontage} + \theta_2 \times \text{depth}$

- Area
  $x = \text{frontage} \times \text{depth}$

- $h_\theta(x) = \theta_0 + \theta_1 x$

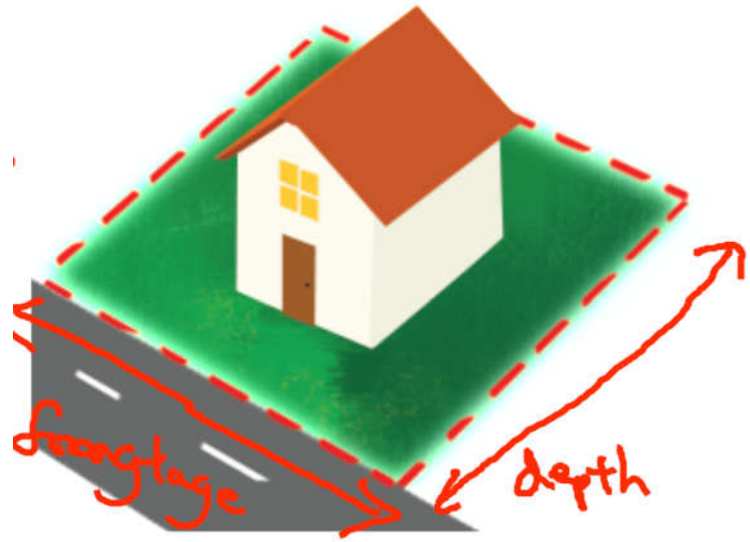

Slide credit: Andrew Ng

Polynomial regression

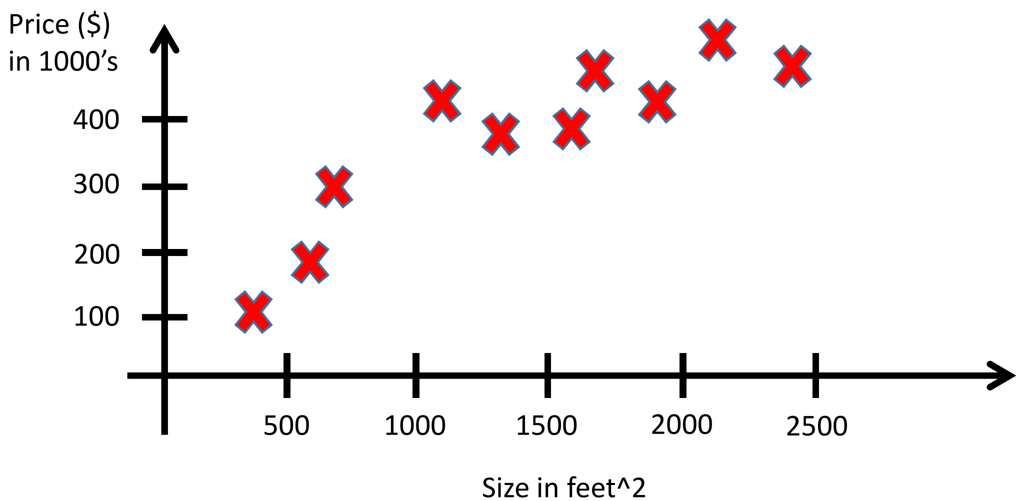

x_1 = (size)
x_2 = (size)^2
x_3 = (size)^3

- $h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$
  $= \theta_0 + \theta_1(size) + \theta_2(size)^2 + \theta_3(size)^3$

Slide credit: Andrew Ng

Linear Regression

- Model representation
- Cost function
- Gradient descent
- Features and polynomial regression
- **Normal equation**

(x_0)	[illegible]	[illegible]	[illegible]	Age of home (years) (x_4)	Price ($) in 1000's (y)
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178
	…				…

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

$$y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

$$\theta = (X^\top X)^{-1} X^\top y$$

Slide credit: Andrew Ng

m training examples, n features

Gradient Descent

- Need to choose α
- Need many iterations
- Works well even when n is large

Normal Equation

- No need to choose α
- Don't need to iterate
- Need to compute $(X^\top X)^{-1}$
- Slow if n is very large

Slide credit: Andrew Ng

Things to remember

- **Model representation**

$$h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n = \theta^\top x$$

- **Cost function** $J(\theta) = \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)^2$

- **Gradient descent for linear regression**

  Repeat until convergence $\{\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right) x_j^{(i)}\}$

- **Features and polynomial regression**

  Can combine features; can use different functions to generate features (e.g., polynomial)

- **Normal equation** $\theta = (X^\top X)^{-1} X^\top y$