



Artificial Intelligence

COURSE INSTRUCTOR: MUHAMMAD SAIF UL ISLAM

Outline

- Hill-Climbing Search.
- Simulated Annealing Search.
- Local Beam Search.
- Genetic Algorithms.

Classical Search vs. Local Search


Classical search	Local Search
<ul style="list-style-type: none">• systematic exploration of search space.• Keeps one or more paths in memory.• Records which alternatives have been explored at each point along the path.• The path to the goal is a solution to the problem.	<ul style="list-style-type: none">• In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution.• State space = set of "complete" configurations.• Find configuration satisfying constraints, Find best state according to some objective function $h(s)$. e.g., n-queens, $h(s)$ = number of attacking queens. In such cases, we can use Local Search Algorithms.

Local Search Algorithms

- Local Search Algorithms keep a **single "current"** state and move to neighboring states to try to **improve** it.
- Solution **path** needs not to be maintained.
- Hence, the search is “**local**”.
- Local search **suitable** for problems in which path is not important; the goal state itself is the solution.
- It is an **optimization** search

Example: n-queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal.
- In the 8-queens problem, what matters is the final configuration of queens, not the order in which they are added.



Local Search: Key Idea

Key idea:


1. Select (**random**) initial state (generate an initial guess).
2. Make local modification to improve current state
(evaluate current state and move to other states).
3. Repeat Step 2 until goal state found (or out of time).

Local Search: Key Idea

Advantages	Drawback:
<ul style="list-style-type: none">• Use very little memory - usually a constant amount.• Can often find reasonable solutions in large or infinite state spaces (e.g., continuous). For which systematic search is unsuitable.	<ul style="list-style-type: none">• Local Search can get stuck in local maxima and not find the optimal solution.

State-Space Landscape

- A state space landscape: is a graph of states associated with their costs.
- State-space landscape
 - Location (defined by state)
 - Elevation (defined by the value of the heuristic cost function or objective function)
 - If elevation = cost, aim to find the lowest valley (**a global minimum**)
 - If elevation = objective function, find the highest peak (**a global maximum**)
 - *A complete local search algorithm always find a goal if one exists*
 - An **optimal** algorithm always find a **global minimum/maximum**




Local and Global Optima

- Global optimum
 - A solution that is better than all other solutions
 - Or no worse than any other solution
- Local optimum
 - A solution which is better than *nearby* solutions
 - A local optimum is not necessarily a global one

Global /Local(max/min)

- A local max/min is over a small area.
 - For instance, if a point is lower than the next nearest point on the left & right then it's a local min.
- There can be many local maxes and mins over an entire graph.
- A global max/min is the highest/lowest point on the entire graph.
- There can only be ONE global max and/or min on a graph and there may not be one at all.


Global /Local(max/min)



Hill-Climbing Search

Hill-Climbing Search

- **Main Idea:** Keep a single current node and move to a neighboring state to improve it.
- Uses a loop that continuously moves in the direction of increasing value (**uphill**):
 - Choose the best successor, choose **randomly** if there is more than one.
 - Terminate when a peak is reached where no neighbor has a higher value.
- It is also called greedy local search, steepest ascent/descent.



Hill-Climbing Search

- “Like climbing Everest in thick fog with amnesia”
- Only record the state and its evaluation instead of maintaining a search tree

function HILL-CLIMBING (*problem*) **returns** a state that is a local maximum

inputs: *problem*, a problem

local variables: *current*, a node
 neighbor, a node

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])


loop do

neighbor \leftarrow a highest-valued successor of *current*

if VALUE[*neighbor*] \leq VALUE[*current*] **then return** STATE[*current*]

current \leftarrow *neighbor*

Hill-Climbing Search




Hill-Climbing Search

Local maxima: a local maximum is a **peak that is higher than each of its neighboring states, but lower than the global maximum.** Hill-climbing algorithms that reach the vicinity of a local maximum will be drawn upwards towards the peak, but will then be stuck with nowhere else to go.

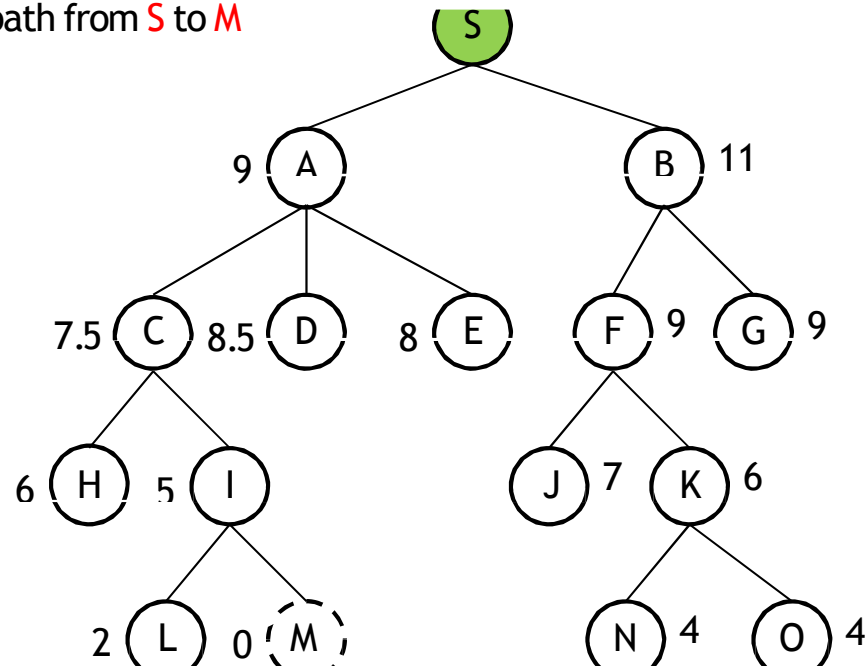
Plateaux: a plateau **is an area of the state space landscape where the evaluation function is flat.** It can be a flat local maximum, from which no uphill exit exists, or a **shoulder**, from which it is possible to make progress.

Ridges: Ridges result in a sequence of local maxima that is very difficult for greedy algorithms to navigate. **(the search direction is not towards the top but the side)**

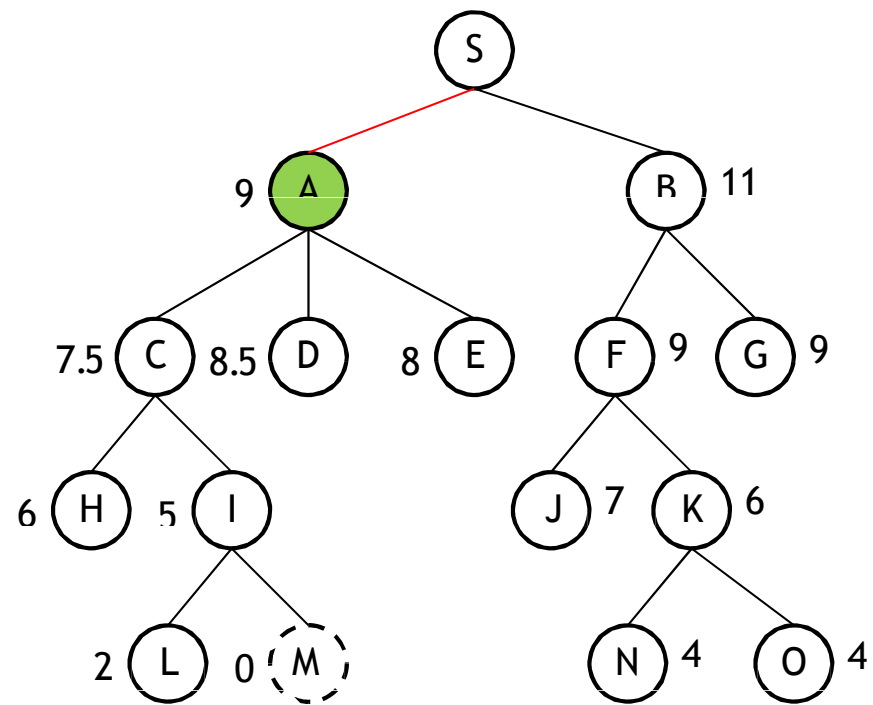


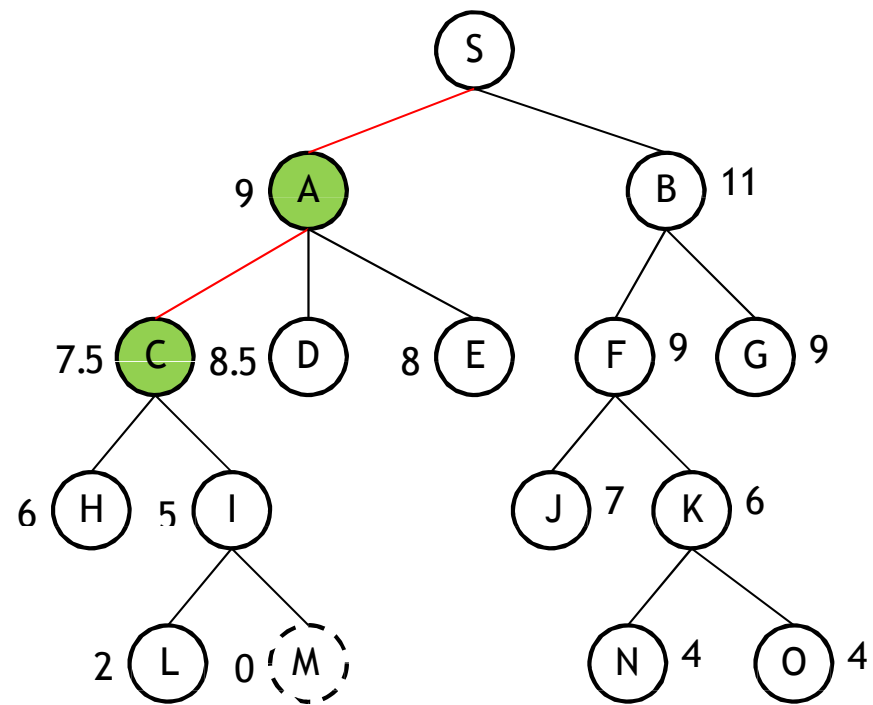
Hill-Climbing Search Example

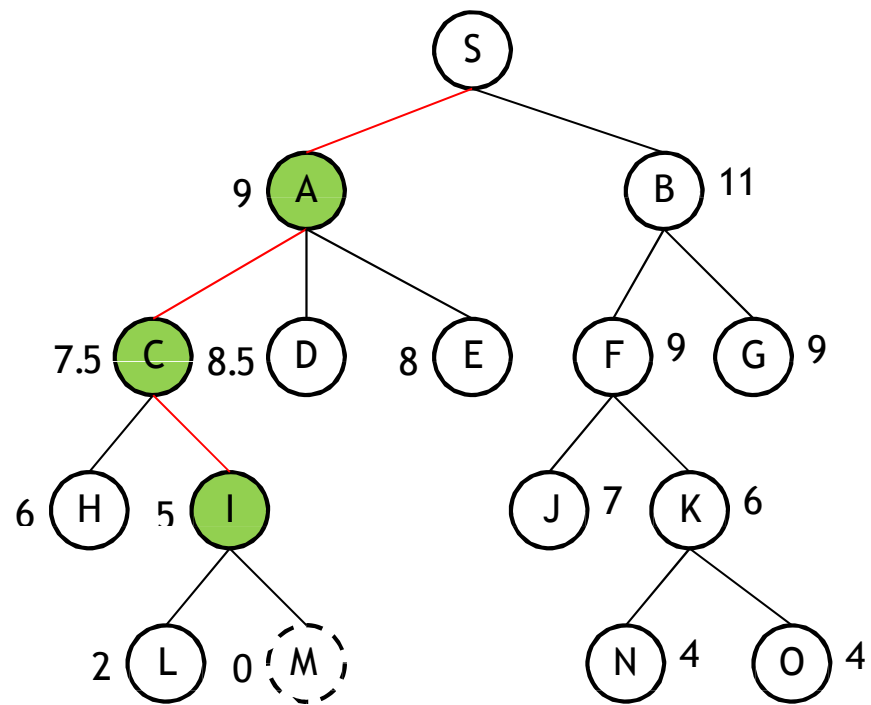
our aim is to find a path from **S** to **M**

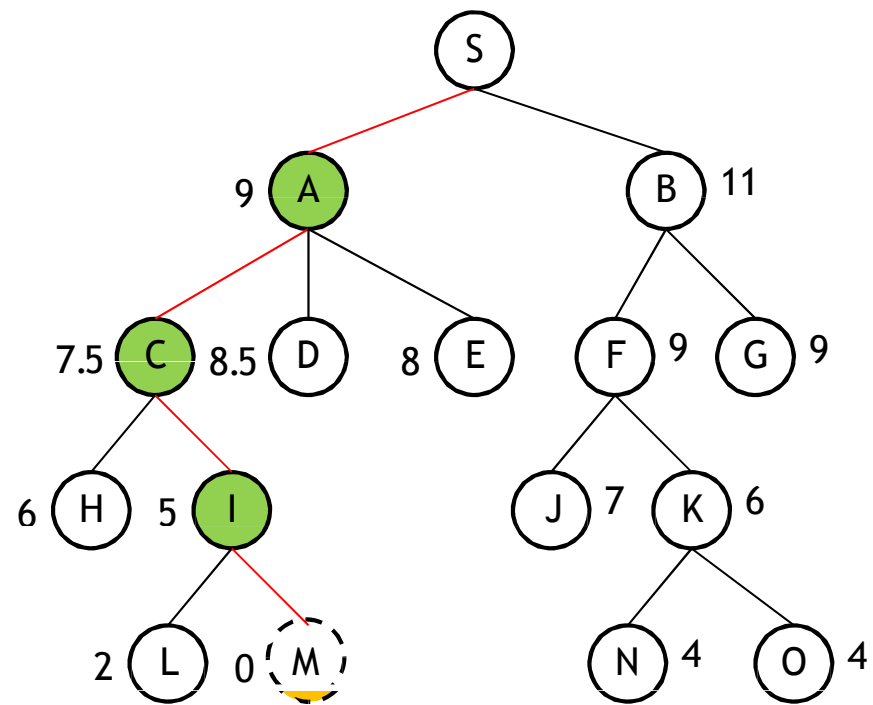


associate
heuristics with
every node, that
is the straight
line distance
from the path
terminating city
to the goal city








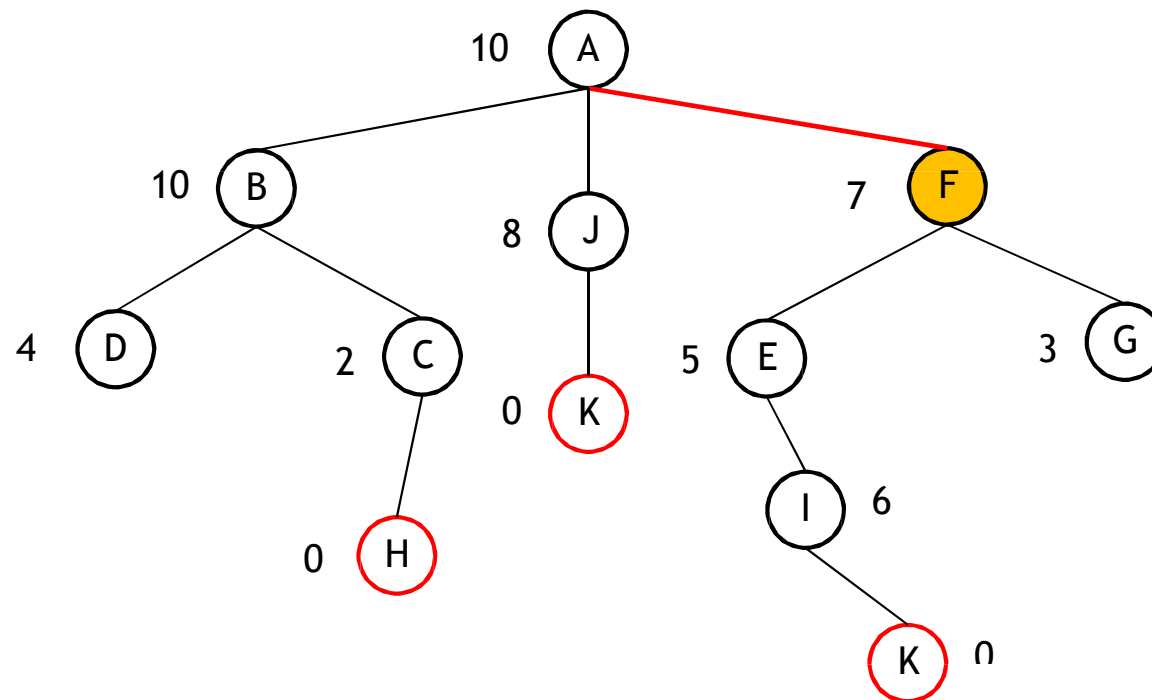


Hill-Climbing Search Example


From **A** find a solution where
H and **K** are finalstates



Local Maximum




Local Minimum



Hill climbing is sometimes called **greedy local search** because it grabs a good neighbor state without thinking ahead about where to go next.

Local Maximum, Local Minimum



Alternative hill climbing

- Stochastic hill climbing
 - chooses at **random** from among the **uphill moves** (neighbors); the **probability** of selection can vary with the steepness of the uphill move.
 - This usually converges more slowly than the steepest ascent, but in some state landscapes, it finds better solutions.
- First-choice hill climbing
 - implements stochastic hill-climbing by generating successors **randomly** until one is generated that is **better** than the **current state**.
 - This is a **good strategy** when a state has **many** (e.g., thousands) of **successors**.
- Random-restart hill climbing
 - adopts the well-known adage(proverb), "If at first, you don't succeed, try, try again." It conducts a series of hill-climbing searches from the **randomly generated initial state**, stopping when a goal is found. It is complete with **probability** approaching **1**, for the trivial reason that it will eventually generate a goal state as the initial state. (It iteratively does hill-climbing, each time with a random initial condition)