

# Parallel and Distributed Computing

## CS3006 (BCS-6C/6D)

### Lecture 01

Instructor: Dr. Syed Mohammad Irteza

Assistant Professor, Department of Computer Science, FAST

24 January, 2023

# Administrative Information

- Office: L-109, Upper Floor, Library Building
- Email: [m.irteza@nu.edu.pk](mailto:m.irteza@nu.edu.pk) & [mohammad.irteza@lhr.nu.edu.pk](mailto:mohammad.irteza@lhr.nu.edu.pk)
- Office Hours: (most probably Monday/Wednesday)
- Website: Google Classroom will be announced soon, iA

# My Academic & Research Background

- BS, MS and PhD in Computer Science from LUMS
  - 2002, 2005, 2018
- PhD Thesis: *Resilient Network Load Balancing for Datacenters*
  - Advisor – Dr. Ihsan Ayyub Qazi
- Google Scholar Page:
  - <https://scholar.google.com/citations?hl=en&user=wHazKsgAAAAJ>
- Main Research Interests:
  - Networking for Datacenters: network layer and transport layer protocols
  - Software Defined Networking

# Teaching Experience

- **FAST** – February 2021 to date
  - Parallel & Dist. Computing, Comp. Networks, Programming Fundamentals, Adv. Statistics
- **UMT** – October 2019 to February 2021
  - Grad-level courses → Advanced Networks, Advanced Computer Architecture
  - Programming Fundamentals, Artificial Intelligence, Computer Networks, Cloud Computing
- **GIFT University**, Gujranwala – April 2019 to August 2019
  - Software Requirement Engineering
- **LUMS** – September 2014 to May 2017
  - Co-instructor/grader for Intro to Programming (C++) and Machine Learning
- **Air University**, Islamabad – January 2010 to August 2012
  - Database Systems, Advanced Database Systems, Digital Logic Design, Computer Architecture, Distributed Systems, Number Theory
- **FAST**, Islamabad – Summer 2008
  - Programming for Engineers-II (C++)

# Software Industry Experience

- ***Prosol Technologies***, Islamabad (partially now [Ciklum](#))
  - Software Consultant (Sep 2008 – Nov 2008)
  - Senior Software Engineer (Feb 2006 – Feb 2008)
  - Software Engineer (Apr 2004 – June 2004)
- ***Diyatech Pakistan*** ([Alachisoft](#)), Islamabad
  - Software Developer (Apr 2003 – Apr 2004)
- [InvestCorp](#), Bahrain
  - Junior Software Developer (June – July 2000)

# Interest in Distributed Computing?

- Undergraduate
  - Distributed Systems (Dr. Salman Iqbal)
  - Data Communications (Dr. Syed Ijlal Shah)
  - Network Programming in Java (Dr. Humaira Kamal)
  - FYP: Enhanced Java Parallel Virtual Machine (Dr. Humaira Kamal)
- MS →
  - High Performance Computing (Dr. Asim Karim)
  - Distributed Software System Development (Mr. Umair Javed, CEO-tkxel)
- PhD →
  - Distributed Systems (Dr. Basit Shafiq)
  - Topics in Internet Research (Dr. Ihsan Ayyub Qazi/Dr. Zartash Uzmi)

# Classroom Etiquette

- Please come on time
- Talking among each other is not acceptable, *while I am teaching*
- Leaving the class to attend a phone call *is not appreciated*
- Quizzes will in general be *unannounced*
  - They can be held at the start or end of class
- Cases of *plagiarism* (copying of other people's work) will lead to marks and/or grade *reductions*

# Grading Policy – Tentative (*may be changed*)

- Quizzes & Assignments → 15% + 15%
  - If we have 7 or more quizzes, we will choose the best 5 or 6
  - All assignments will count to your grade
- Midterm I and Midterm II → 30%
- Final Exam → 40%
  - Comprehensive exam (all course contents included)



# Textbooks

- *Introduction to Parallel Computing* by Ananth Grama and Anshul Gupta.
- *Distributed Systems: Concept and Design* by George Coulouris, Gordon Blair
- *Using OpenMP: Portable Shared Memory Parallel Programming* by Barbara Chapman, Gabriele Jost, Ruud van der Pas.

# Reference Books

- *Distributed Systems: Principles and Paradigms*, A. S. Tanenbaum and M. V. Steen, Prentice Hall, 2nd Edition, 2007.
- *Distributed and Cloud Computing: Clusters, Grids, Clouds, and the Future Internet*, K Hwang, J Dongarra and GC. C. Fox, Elsevier, 1st Ed.

# Course Objectives

- To understand the fundamental concepts of *parallel and distributed computing*
- The design and analysis of *parallel algorithms*
- Analyzing different problems and then *developing parallel programming solutions* for those *problems*
- Study the challenges of *Parallel and Distributed Systems* and how to cope with them

# Course Schedule

Week	Topic
01	Introduction to parallel and distributed systems; Motivating parallelism; Amdahl's Law
02	Flynn's Taxonomy, Multithreading
03	Shared Memory Architecture
04	Principles of parallel algorithm design
05	Basic Communication Operations
	Midterm I
06	Programming Shared Address Space Platforms using POSIX Thread API and OpenMP
07	Decompositions techniques, Shared memory programming with OpenMP + <b><i>Project Proposals</i></b>
08	Parallel programming with OpenMP
09	Introduction to Distributed Systems
10	Types of Distributed Systems + <b><i>Project Phase 1</i></b>

# Course Schedule

Week	Topic
	Midterm II
11	Programming Distributed machines using Message passing interface (MPI)
12	Collective Communication and Computation Operations
13	Fault Tolerance Techniques
14	Project Presentations + <b><i>Project Phase 2</i></b>
15	Project Presentations
	Final Exam

# The Changing Nature of Applications

- The *scale* of the *user-base* for many *popular user-facing services* is so *large*, that *traditional models* for *hosting* and *deploying* applications will not work

# Cloud enables highly leveraged services, at scale

- **Facebook:**

- ~*2.8 billion MAU* (monthly active users, 2021)
- ~*1.8 billion DAU* (including FB, WhatsApp, Insta, Messenger)
- Generates *no content* itself
- Disrupts *media companies*



- **Uber:**

- Ride sharing company → *93m customers, 3.5m drivers* ([2021](#))
- Owns *no vehicle*
- Disrupts *multiple markets* (\$26.6b gross bookings, 2020)
  - *Taxi services*
  - *Vehicle ownership*



# Services live in Clouds

- Infrastructure

- Datacenters (DCs)
- Clusters/pods
- Rows/racks
- Servers/switches

- Deployment

- Public
- Private
- Hybrid






# Discussion for today

- Motivating Parallelism
- Computing v. Systems
- Parallel and Distributed Computing
- Practical Applications of P & D Computing

# Motivating Parallelism

- *Uniprocessor* are *fast* but:
  - Some problems require *too much computation*
  - Some problems use *too much data*
  - Some problems have *too many parameters to explore*
- For example:
  - Weather simulations, gaming, web servers, code breaking



# Motivating Parallelism


- Developing *parallel hardware and software* has traditionally been *time and effort intensive*.
- If one is to view this in the context of *rapidly improving uniprocessor speeds*, one is tempted to question the *need for parallel computing*.
- Latest trends in *hardware design* indicate that *uni-processors* may not be able to sustain the *rate of realizable performance increments* in the future .
- This is the result of a number of *fundamental physical and computational limitations*.
- The emergence of *standardized parallel programming environments, libraries,* and *hardware* have *significantly reduced time to develop (parallel) solutions*.

# Motivating Parallelism – Moore's Law

- Proposed by *Gordon E. Moore* in *1965* and revised in *1975*.
- It states that [*Simplified Version*]:
  - “*Processing speeds, or overall processing power for computers will double every 18 months.*”
- A more technically correct interpretation:
  - “The *number of transistors on an affordable CPU* would double every two years [*18 months*].”

# Moore's Law

- *Number of transistors* incorporated in a chip will *approximately double every two years*



# Moore's Law


- More computational power *implicitly means more transistors*.
- Let's have a look at the empirical data from 1970 to 2009
  - In the 1970's (i.e., from 1970 to 1979), processor speeds ranged from 740 KHz to 8 Mhz. The difference shows that both the interpretations are correct.
  - From 2000 to 2009, Speeds ranged from 1.3 GHz to 2.8 GHz.
  - Speed difference is too low but, *number of integrated transistors ranged from 37.5 million to 904 million*.

# Moore's Law

- Why doubling the transistors does not double the speed?
  - The answer is increase in number of transistor per processor is due to multi-core CPU's.
  - It means, to follow Moore's law, companies had to:
    - Introduce ULSI (ultra large-scale integrations)
    - And multi-core processing era.
- Will Moore's law hold forever?
  - Adding multiple cores on single chip causes heat issues.
  - Furthermore, increasing the number of cores, may not be able to increase speeds [Due to inter-process interactions].
  - Moreover, transistors would eventually reach the limits of miniaturization at atomic levels

# Moore's Law

- So, we must look for efficient parallel software solutions to fulfill our future computational needs.
- As stated earlier, number of cores on a single chip also have some restrictions.
- Solution(s)?
  - Need to find more scalable distributed and hybrid solutions





# Motivating Parallelism

## *The Memory/Disk Speed Argument*

- While clock rates of high-end processors have increased at roughly 40% per year over the past decade, DRAM access times have only improved at the rate of roughly 10% per year over this interval.
- This mismatch in speeds causes significant performance bottlenecks.
- Parallel platforms provide increased bandwidth to the memory system.
- Parallel platforms also provide higher aggregate caches.
- Some of the fastest growing applications of parallel computing utilize not their raw computational speed, rather their ability to pump data to memory and disk faster.

# Motivating Parallelism

## *The Data Communication Argument*

- As the network evolves, the vision of the Internet as one large computing platform has emerged.
- In many applications like databases and data mining problems, the volume of data is such that they cannot be moved.
- Any analyses on this data must be performed over the network using parallel techniques

# Computing v. Systems

## *Distributed Systems*

- A collection of autonomous computers, connected through a network and distribution middleware.
  - This enables computers to coordinate their activities and to share the resources of the system.
  - The system is usually perceived as a single, integrated computing facility.
  - Mostly concerned with the hardware-based accelerations

## *Distributed Computing*

- A specific use of distributed systems, to split a large and complex processing into subparts and execute them in parallel, to increase the productivity.
  - Computing mainly concerned with software-based accelerations (i.e., designing and implementing algorithms)

# Sources

- Slides of Dr. Rana Asif Rahman, FAST