

# Parallel and Distributed Computing

## CS3006 (BCS-6C/6D)

### Lecture 21

Instructor: Dr. Syed Mohammad Irteza

Assistant Professor, Department of Computer Science, FAST

13 April, 2023

# Previous Lecture

- MPI
  - MPI\_Probe and MPI\_Get\_count
  - MPI\_Barrier
  - MPI\_Bcast
  - MPI\_Reduce (different operations)
  - MPI\_Allreduce
  - MPI\_Scan (different operations)
  - MPI\_Gather and MPI\_Scatter
  - SPMD Model
  - MPI\_Alltoall

# Sorting in the Parallel Era.....

- Can we efficiently apply a Bubble-Sort type of sorting algorithm when the individual values are dispersed across different machines (processes)?

# Sorting - Overview

- One of the most commonly used and well-studied Algorithms.
- Sorting can be *comparison-based* or *non-comparison-based*.
- The fundamental operation of comparison-based sorting is *compare-exchange*.
- The lower bound on any comparison-based sort of  $n$  numbers is  $\Theta(n \log n)$  .
- Let's explore a comparison-based sorting algorithm.

# Sorting – Basics

What is a parallel sorted sequence?


→ Where are the input and output lists stored?

## Answers:

- We assume that the input and output lists are distributed.
- The sorted list is partitioned with the property that each partitioned list is sorted and each element in processor  $P_i$ 's list is less than that in  $P_j$ 's list if  $i < j$ .

# Sorting: Parallel Compare Exchange Operation

- A parallel compare-exchange operation. Processes  $P_i$  and  $P_j$  send their elements to each other. Process  $P_i$  keeps  $\min\{a_i, a_j\}$ , and  $P_j$  keeps  $\max\{a_i, a_j\}$ .



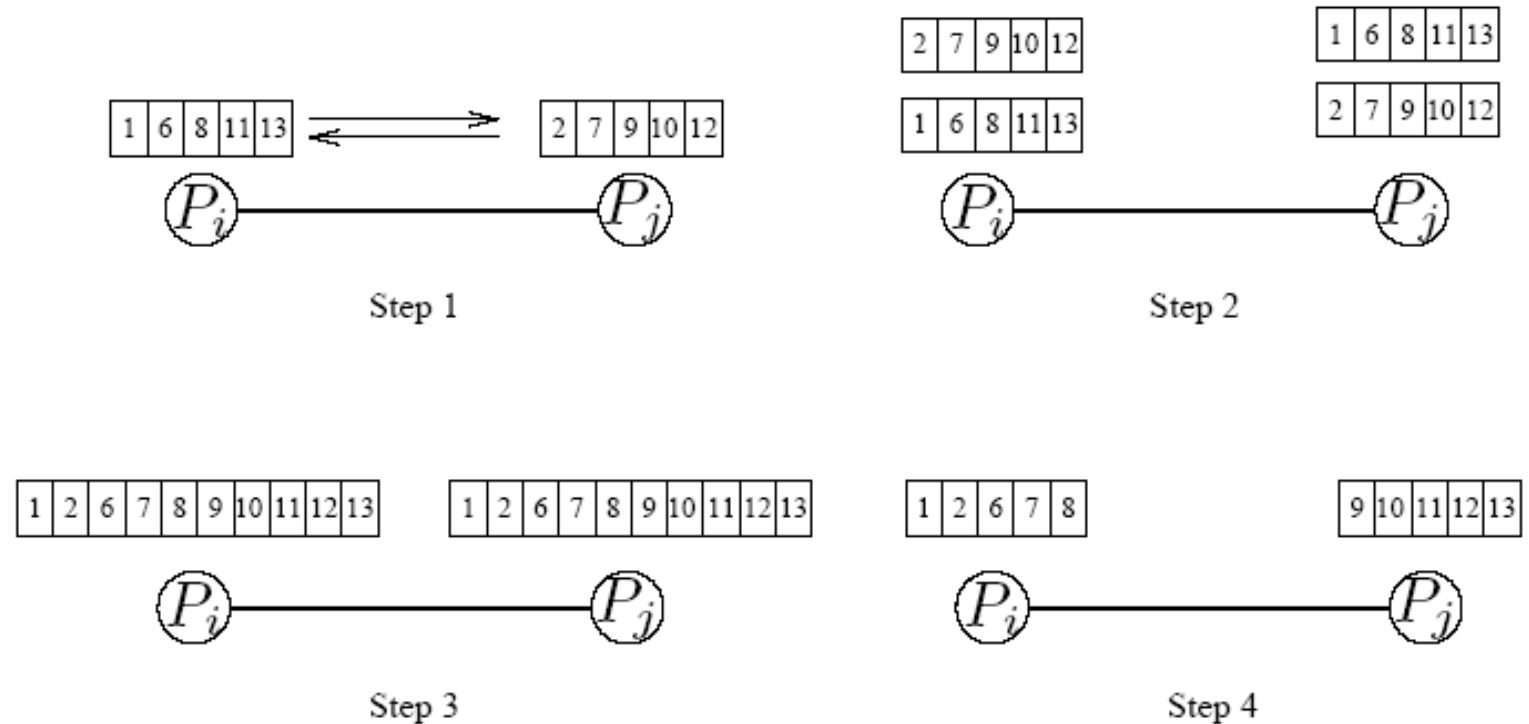
# Sorting: Parallel Compare Exchange Operation

## [cost estimation]

- If each processor has one element, the compare exchange operation stores the smaller element at the processor with smaller id. This can be done in  $t_s + t_w$  time.
- If we have more than one element per processor, we call this operation a compare split. Assume each of two processors have  $n/p$  elements.
- After the compare-split operation, the smaller  $n/p$  elements are at processor  $P_i$  and the larger  $n/p$  elements at  $P_j$ , where  $i < j$ .
- The time for a compare-split operation is  $(t_s + t_w n/p)$ , assuming that the two partial lists were initially sorted.
  - Note that this time is only accounting communication costs. Computation and memory complexities are separate things.

# Sorting: Parallel Compare Exchange

- A compare-split operation. Each process sends its block of size  $n/p$  to the other process.
- Each process merges the received block with its own block and retains only the appropriate half of the merged block.
- In this example, process  $P_i$  retains the smaller elements and process  $P_j$  retains the larger elements.



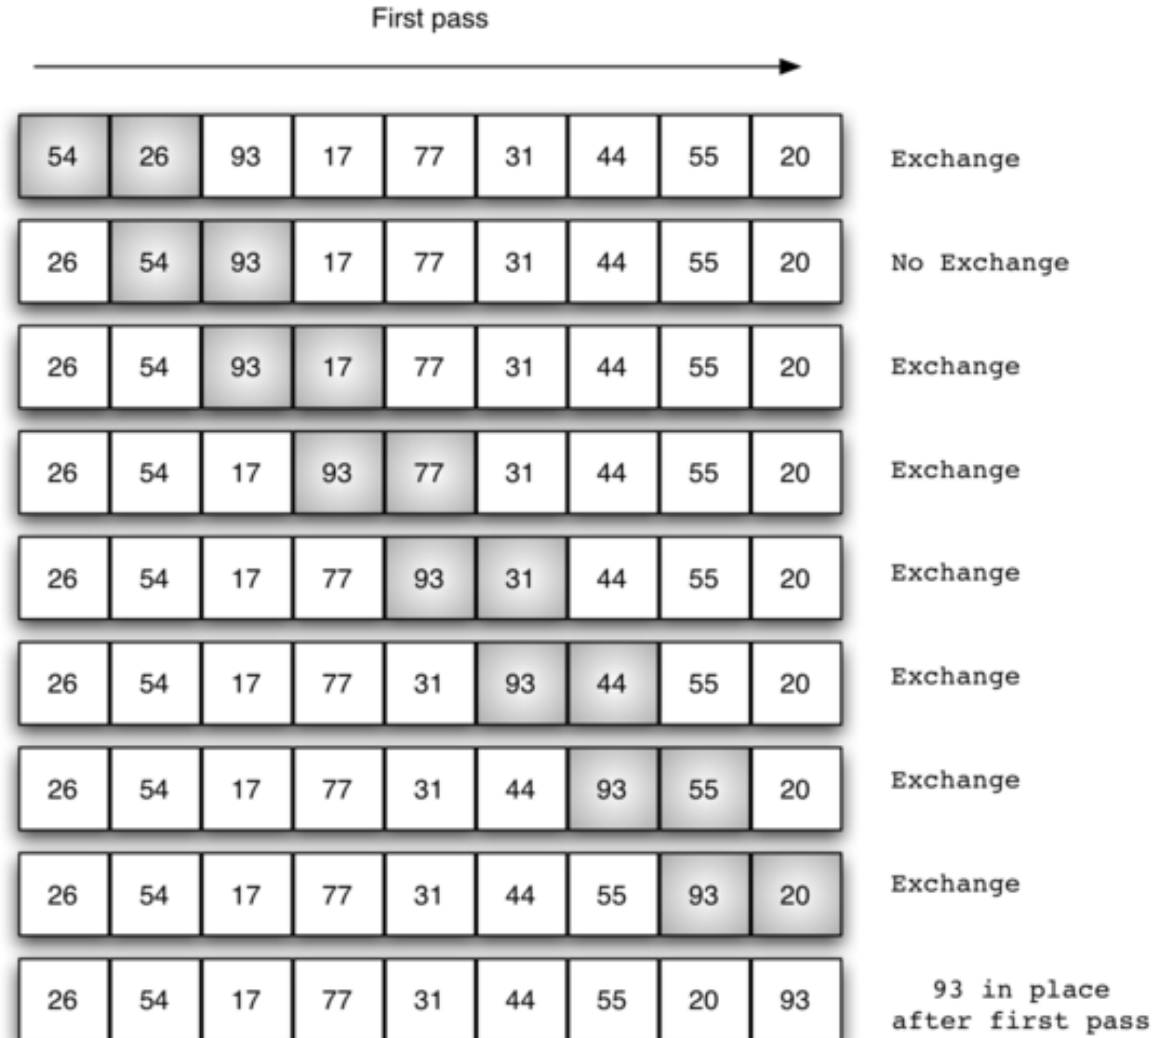


# Bubble Sort and its Variant

- The sequential bubble sort algorithm compares and exchanges adjacent elements in the sequence to be sorted:

```
1.      procedure BUBBLE_SORT( $n$ )  
2.      begin  
3.          for  $i := n - 1$  downto 1 do  
4.              for  $j := 1$  to  $i$  do  
5.                  compare-exchange( $a_j, a_{j+1}$ );  
6.      end BUBBLE_SORT
```

# Visualization (first pass, seq. BubbleSort)




# Bubble Sort and its Variant

- The complexity of bubble sort is  $\Theta(n^2)$ .
- Bubble sort is difficult to parallelize since the algorithm has no concurrency.
- A simple variant, though, uncovers the possible concurrency.

# Bubble Sort [Odd-Even Transposition]

- Sorting  $n = 8$  elements, using the odd-even transposition sort algorithm.
- During each phase, at most 8 elements are compared.
- **[This is according to the sequential algorithm]**



# Bubble Sort [Odd-Even Transposition]


- Sequential odd-even sort algorithm

```
1.  procedure ODD-EVEN( $n$ )
2.  begin
3.      for  $i := 1$  to  $n$  do
4.      begin
5.          if  $i$  is odd then
6.              for  $j := 0$  to  $n/2 - 1$  do
7.                  compare-exchange( $a_{2j+1}, a_{2j+2}$ );
8.          if  $i$  is even then
9.              for  $j := 1$  to  $n/2 - 1$  do
10.                 compare-exchange( $a_{2j}, a_{2j+1}$ );
11.      end for
12.  end ODD-EVEN
```

# Odd-Even Sort (Seq. Complexity)

- After  $n$  phases of odd-even exchanges, the sequence is sorted.
- Each phase of the algorithm (either odd or even) requires  $\Theta(n)$  comparisons.
- Serial complexity is  $\Theta(n^2)$ .

# Parallel Odd-Even Sort



Parallel time complexity:  $T_{par} = O(n)$  (for  $P=n$ )

Source: <https://www.slideshare.net/richakumari37266/parallel-sorting-algorithm>

# Parallel Odd-Even Sort

- **Algorithm Through Observations:**

1. There are total **P** phases/steps. Where P is number of processes
2. **For even phases**
  - i. If 'myrank' is even  $\rightarrow$  Communication partner is ('myrank'+1)
  - ii. If 'myrank' is odd  $\rightarrow$  Communication partner is ('myrank' - 1)
3. **For odd phases:**
  - i. If 'myrank' is even  $\rightarrow$  Communication partner is ('myrank' - 1)
  - ii. If 'myrank' is odd  $\rightarrow$  Communication partner is ('myrank'+1)
4. Communication partners remain constant
5. If 'myrank' is less-than the partner, then keep lower values in compare-split-operation



# Parallel Odd-Even Sort

## Complexity when $n = P$

- Consider the one item per processor case.
- There are  $P$  iterations, in each iteration, each processor does one compare-exchange.
- The parallel run time of this formulation is  $\Theta(n)$ .
- Parallel run time means computation performed by each of the processors in parallel.

# Parallel Odd-Even Sort

## Complexity when $n > P$

- Consider a block of  $n/p$  elements per processor.
- The first step is a local sort.
- In each subsequent step, the compare exchange operation is replaced by the compare split operation.
- The parallel run time of the formulation is:

$$T_P = \overbrace{\Theta\left(\frac{n}{p} \log \frac{n}{p}\right)}^{\text{local sort}} + \overbrace{\Theta(n)}^{\text{comparisons}} + \overbrace{\Theta(n)}^{\text{communication}}$$

comm. steps for a single process

```
1  #include <stdlib.h>
2  #include <mpi.h> /* Include MPI's header file */
3
4  main(int argc, char *argv[])
5  {
6      int n;          /* The total number of elements to be
sorted */
7      int npes;       /* The total number of processes */
8      int myrank;     /* The rank of the calling process */
9      int nlocal;     /* The local number of elements, and the
array that stores them */
10     int *elmnts;    /* The array that stores the local
elements */
11     int *relmnts;   /* The array that stores the received
elements */
12     int oddrank;    /* The rank of the process during odd-
phase communication */
13     int evenrank;   /* The rank of the process during even-
phase communication */
14     int *wspace;    /* Working space during the compare-split
operation */
```

```
15     int i;
16     MPI_Status status;
17
18     /* Initialize MPI and get system information */
19     MPI_Init(&argc, &argv);
20     MPI_Comm_size(MPI_COMM_WORLD, &npes);
21     MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
22
23     n = atoi(argv[1]);
24     nlocal = n/npes; /* Compute the number of elements to be
stored locally. */
25
26     /* Allocate memory for the various arrays */
27     elmnts = (int *)malloc(nlocal*sizeof(int));
28     relmnts = (int *)malloc(nlocal*sizeof(int));
29     wspace = (int *)malloc(nlocal*sizeof(int));
```

```
31  /* Fill-in the elmnts array with random elements */
32  srand(myrank);
33  for (i=0; i<nlocal; i++)
34      elmnts[i] = random();
35
36  /* Sort the local elements using the built-in quicksort routine */
37  qsort(elmnts, nlocal, sizeof(int), IncOrder);
38
39  /* Determine the rank of the processors that myrank needs to communicate during */
40  /* the odd and even phases of the algorithm */
41  if (myrank%2 == 0) {
42      oddrank  = myrank-1;
43      evenrank = myrank+1;
44  }
45  else {
46      oddrank  = myrank+1;
47      evenrank = myrank-1;
48  }
```

```
50  /* Set the ranks of the processors at the end of the linear */
51  if (oddrank == -1 || oddrank == npes)
52      oddrank = MPI_PROC_NULL;
53  if (evenrank == -1 || evenrank == npes)
54      evenrank = MPI_PROC_NULL;
55
56  /* Get into the main loop of the odd-even sorting algorithm */
57  for (i=0; i<npes-1; i++) {
58      if (i%2 == 1) /* Odd phase */
59          MPI_Sendrecv(elmnts, nlocal, MPI_INT, oddrank, 1, relmnts,
60                      nlocal, MPI_INT, oddrank, 1, MPI_COMM_WORLD, &status);
61      else /* Even phase */
62          MPI_Sendrecv(elmnts, nlocal, MPI_INT, evenrank, 1, relmnts,
63                      nlocal, MPI_INT, evenrank, 1, MPI_COMM_WORLD, &status);
64
65      CompareSplit(nlocal, elmnts, relmnts, wspace,
66                  myrank < status.MPI_SOURCE);
67  }
68
69  free(elmnts); free(relmnts); free(wspace);
70  MPI_Finalize();
71 }
```

```
73  /* This is the CompareSplit function */
74  CompareSplit(int nlocal, int *elmnts, int *relmnts, int *wspace,
75              int keepsmall)
76  {
77      int i, j, k;
78
79      for (i=0; i<nlocal; i++)
80          wspace[i] = elmnts[i]; /* Copy the elmnts array into the wspace array */
81
82      if (keepsmall) { /* Keep the nlocal smaller elements */
83          for (i=j=k=0; k<nlocal; k++) {
84              if (j == nlocal || (i < nlocal && wspace[i] < relmnts[j]))
85                  elmnts[k] = wspace[i++];
86              else
87                  elmnts[k] = relmnts[j++];
88          }
89      }
90      else { /* Keep the nlocal larger elements */
91          for (i=k=nlocal-1, j=nlocal-1; k>=0; k--) {
92              if (j == 0 || (i >= 0 && wspace[i] >= relmnts[j]))
93                  elmnts[k] = wspace[i--];
94              else
95                  elmnts[k] = relmnts[j--];
96          }
97      }
98  }
```

```
100  /* The IncOrder function that is called by qsort is defined as follows */
101  int IncOrder(const void *e1, const void *e2)
102  {
103      return (*((int *)e1) - *((int *)e2));
104  }
```

6 of 6





# Failure Handling

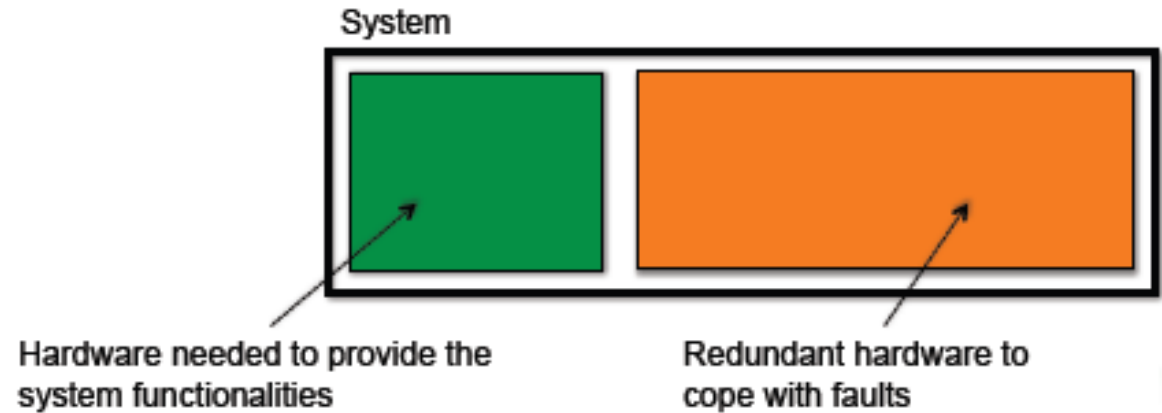
- As the number of components increases, the fault rate also increases
- How to detect failures and make corrections?
- Fault Tolerance
  - fault tolerance or graceful degradation is the property that enables a system to continue operating properly in the event of the failure of some of its components
  - a system continues its intended operation, possibly at a reduced level, rather than failing completely, when some part of the system fails
- Redundancy
  - redundancy is the duplication of critical components or functions of a system with the intention of increasing reliability of the system

# Redundancy design techniques

- Hardware redundancy
- Information redundancy
- Time redundancy


# Hardware redundancy

- The system is implemented using *more hardware than that which is needed* for implementing the system functionalities
- The redundant hardware is used for *dealing with faults*



- Types of Hardware redundancy
  - Passive redundancy
  - Active redundancy
  - Hybrid redundancy

# Passive Redundancy



- The HW needed to implement the system is replicated 3 times:
  - Triple Module Redundancy (TMR)

# Passive Redundancy

- A *majority voter* decides the output to be committed to the user on the basis of the outputs coming from the three domains
- Voter implements the following functionality

```
P_VOTER: process( Y1, Y2, Y3 )  
  Begin  
    if Y1=Y2 then  
      Y = Y1;  
    else if Y1=Y3 then  
      Y = Y1;  
    else if Y2=Y3 then  
      Y = Y2;  
    else "display error";  
    end if;  
  end process;
```

# Passive Redundancy

Used to achieve fault tolerance:

- The voter *stops the propagation of faults* that never reach the outputs: error masking
- In case of *permanent fault of a domain no corrective actions are taken*: no error correction


## Cost:

- Area: >3x
- Time: negligible

# Active Redundancy


Alternative to passive redundancy

- It implements fault detection
- It may also possibly implement:
  - Error localization
  - Error containment
  - Error recovery
- The purpose of fault detection is to assert a *signal (Agree/Disagree)* every time the system output *differs from the expected* for the given input





# Active Redundancy

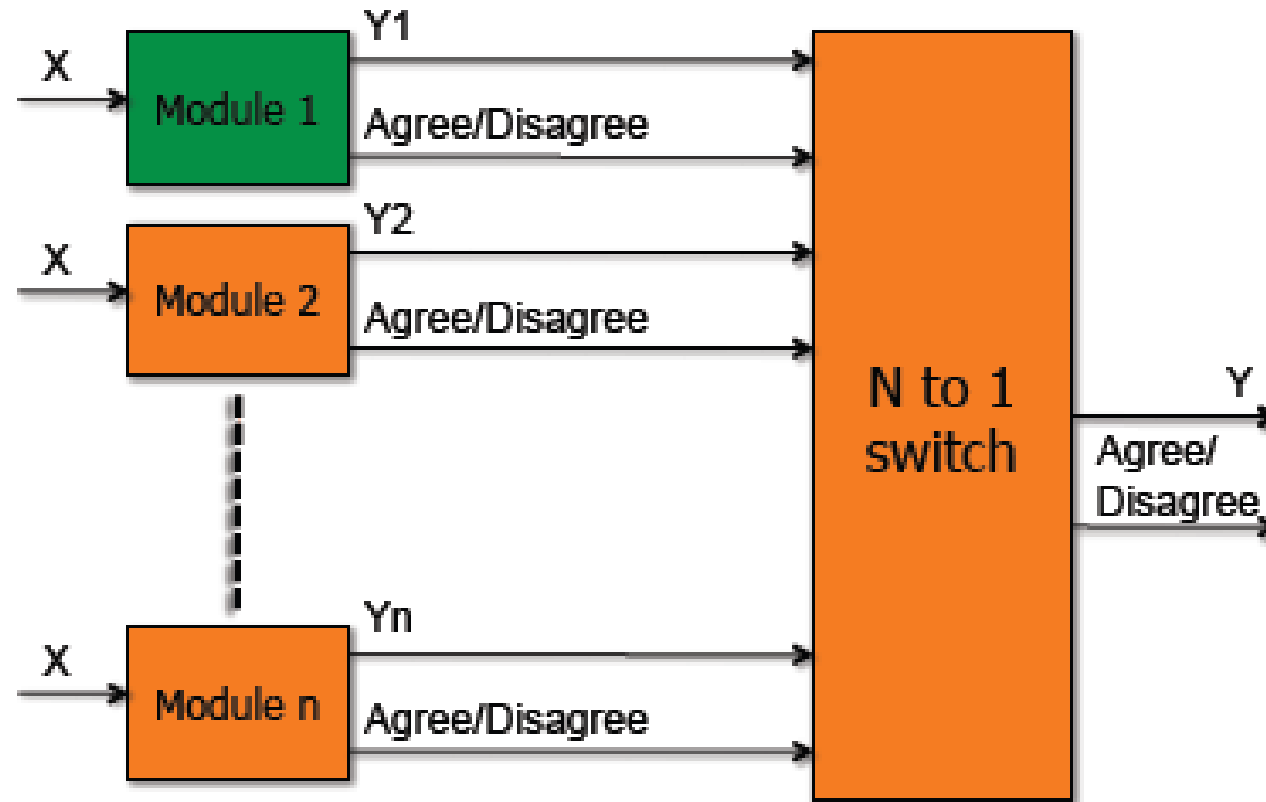


- Duplication with comparison

# Active Redundancy

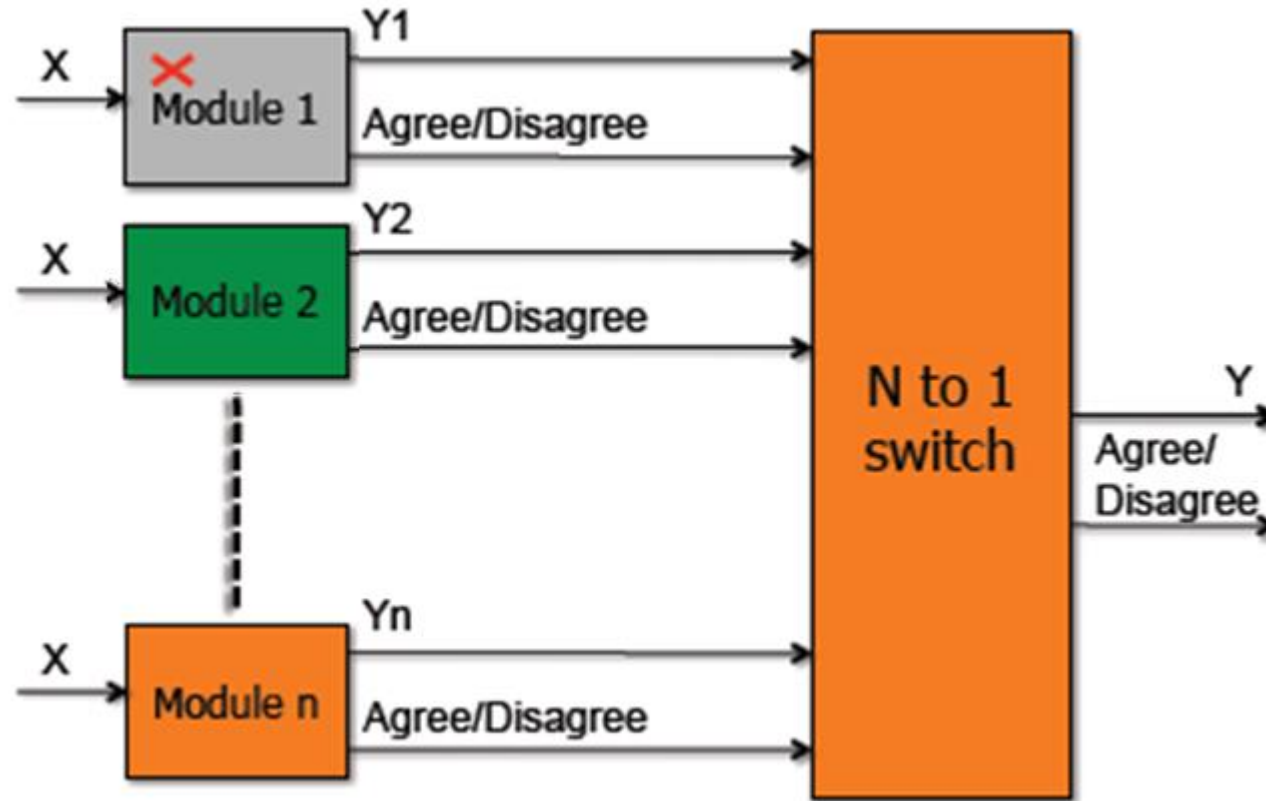
- The purpose of *error localization, containment and recovery* is to reconfigure the system *upon error detection* so that the correct system functionalities are recovered after a certain recovery time
- During recovery the *system is unavailable*:
- The system is not able to respond to inputs
- Maximum allowed recovery time is a function of the application
- Can be implemented using:
  - Hot standby sparing
  - Cold standby sparing

# Hot standby sparing



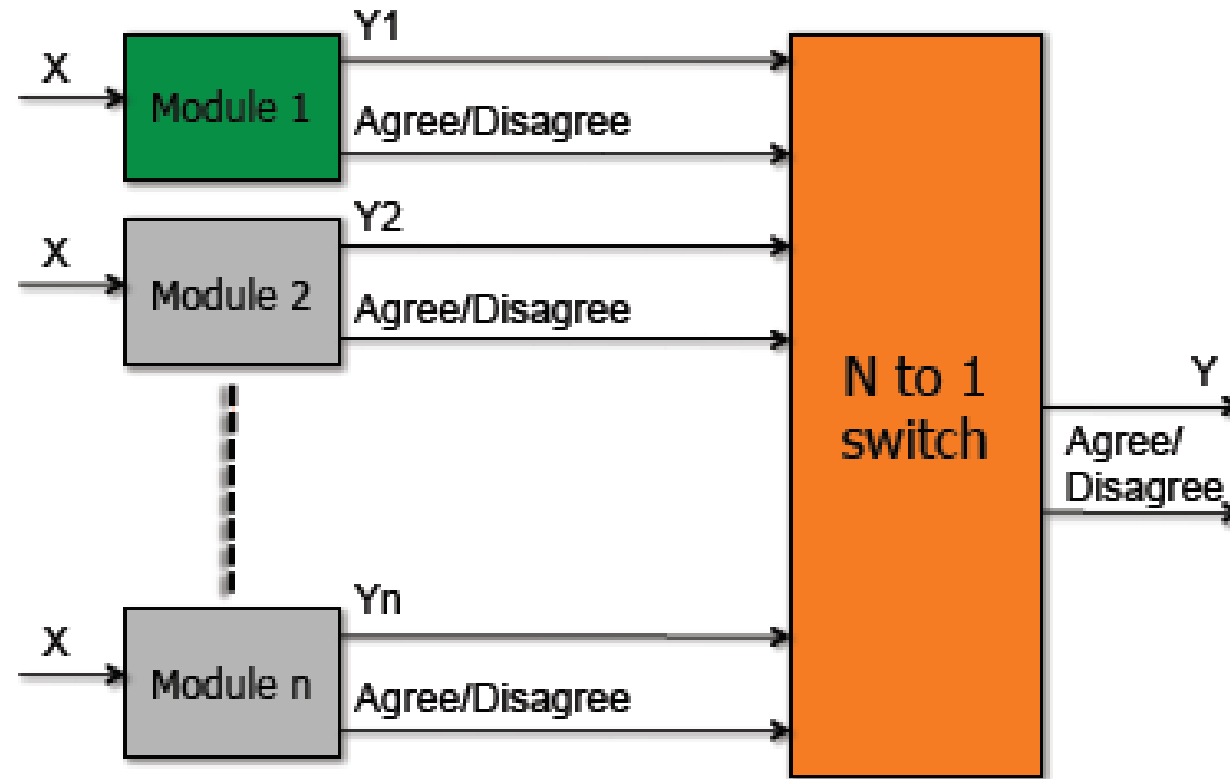
- N modules are used all active at the same time
- One module is the primary module, others are spare modules

# Hot standby sparing



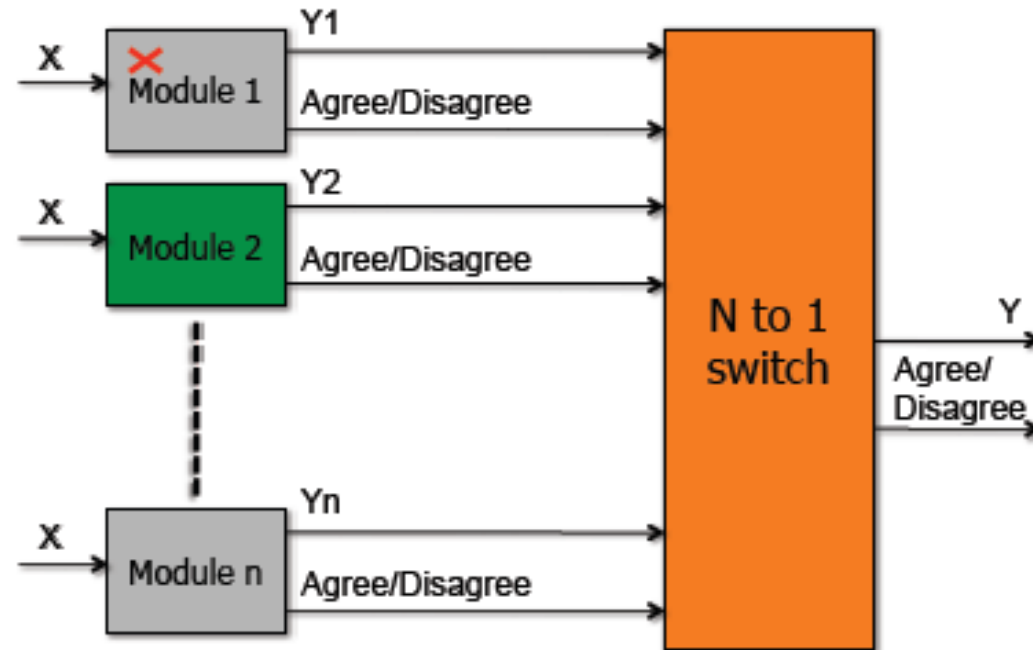
- When an error is detected, the switch sets the faulty primary module to *not-used state*, and it selects a new primary module among the *spares*

# Cold standby sparing



- N modules are used: one active, the *primary module*, the others powered-off, the *spare modules*.

# Cold standby sparing



- When an error is detected, it *powers-off* the *primary module*, selects a new primary module among the *spares*, and then turns it on.

# Dealing with faulty modules

- In hot standby two methods can be used for dealing with the faulty module
- **Permanently evicted:**
  - When detected as faulty, it is powered-off and no longer used
  - The number of spares decreases over time
- **Temporarily evicted:**
  - When detected as faulty, it is set to idle but not powered off
  - The switch monitors the status of the idle module and, in case it is no longer affected by faults, it is promoted to the role of spare module
  - In case of transient faults the number of spares remain constant over time

# Tradeoffs

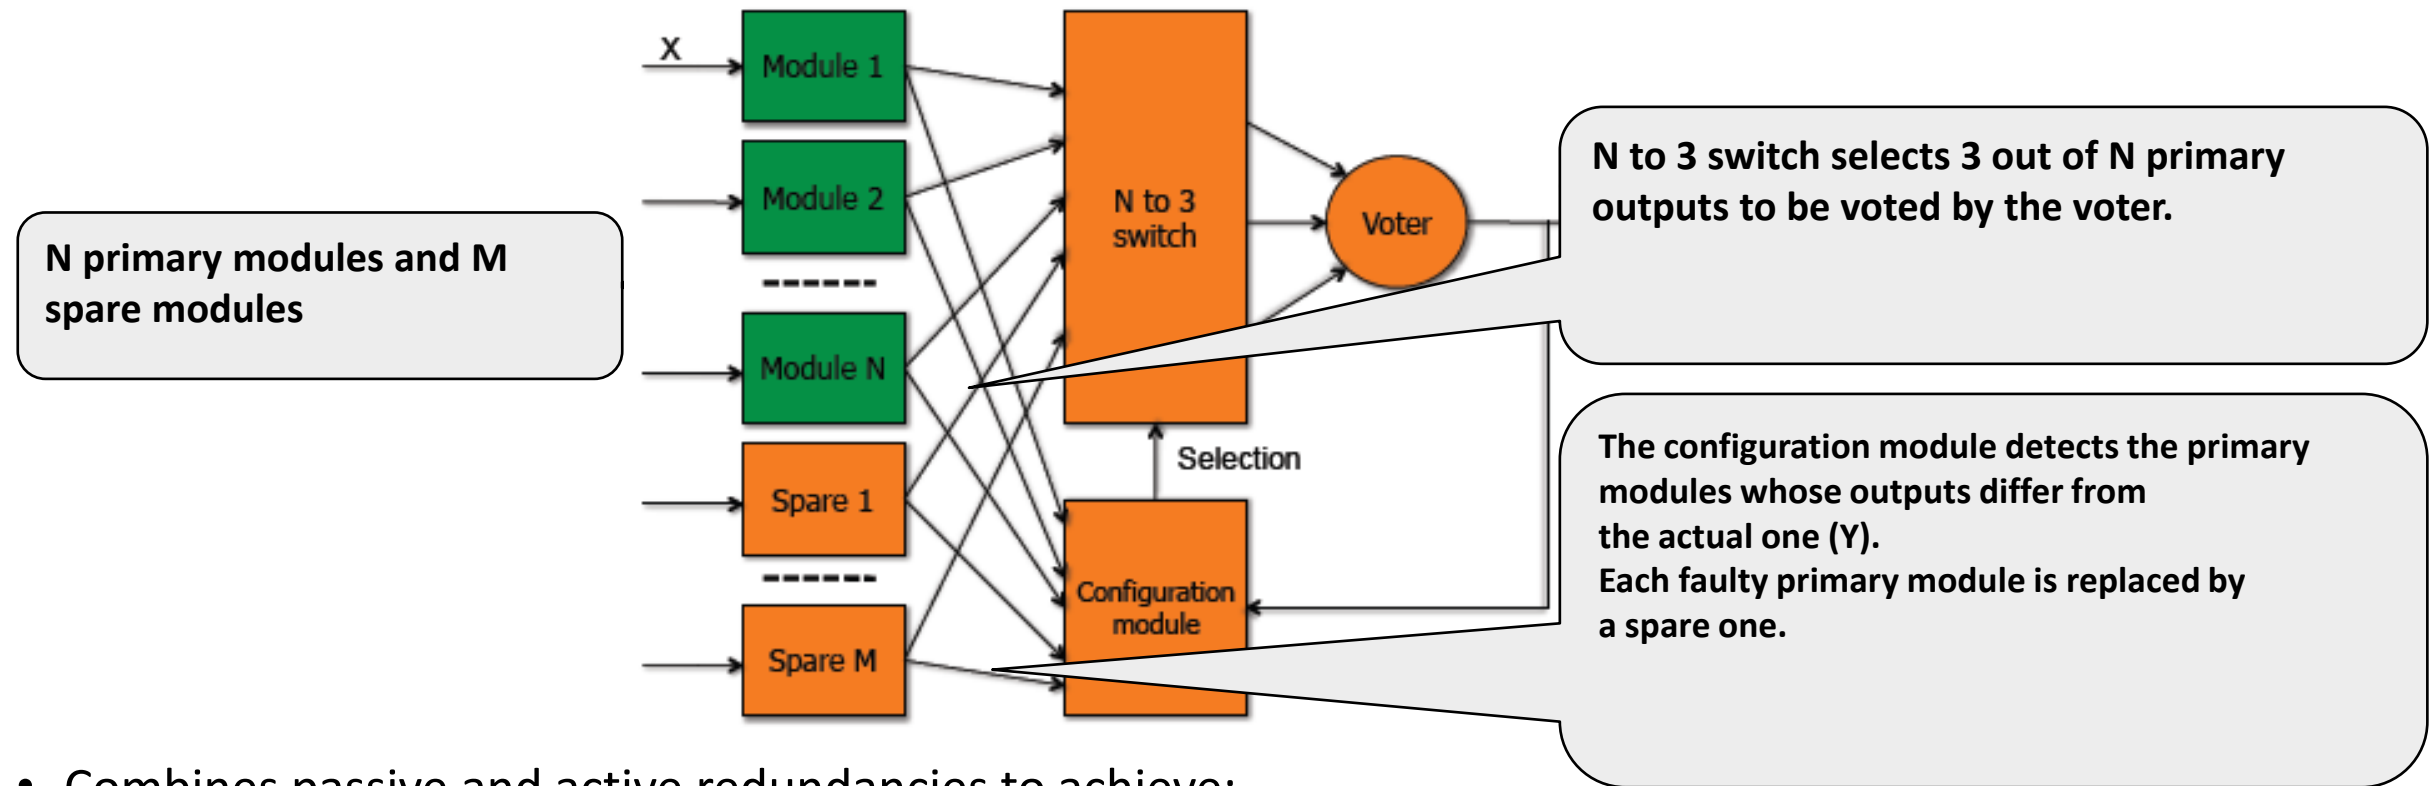
- Active redundancy for error detection only is less expensive than TMR, but does not provide masking
- Standby sparing is more expensive than TMR but it provides error correction:
  - It can survive up to  $N-1$  permanent errors, and provide error detection in case of  $N$  permanent faults
  - TMR is able to survive 1 permanent fault, only

## Cold sparing vs. hot sparing

- **Area:** same
- **Energy:** cold sparing uses  $1/n$  of the energy of hot sparing
- **Availability:** in hot sparing, availability is higher than in cold one



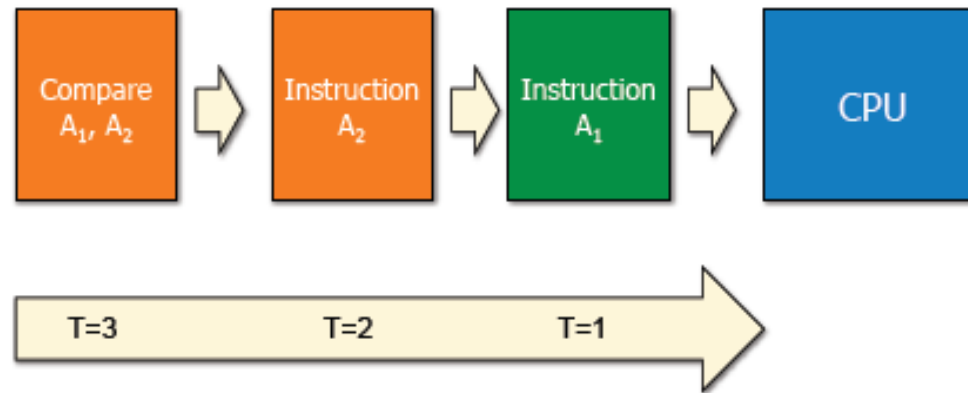
# Hybrid Redundancy



- Combines passive and active redundancies to achieve:
  - Error masking
  - Error correction
- Combining TMR with sparing we can have N modular redundancy with spares

# Time Redundancy


- More time than that needed for processing an input is used
- The additional time is devoted to detect and possibly correct errors occurred during the processing



- The same instruction is executed twice, and a comparison detects the occurrence of errors during the computations

# Time Redundancy

- Error Correction:



- Cost:

- Area: almost negligible
- Time: ~3x for detection, ~4x for correction

# Information Redundancy

- More data than those needed by the application are stored



- The *redundant data* added to the original data is used to detect and possibly *correct errors affecting the original data*
- The redundant data is a function of the original one

# Purpose of information redundancy

- Vulnerable components

- Channels
- Processes (clients, servers)

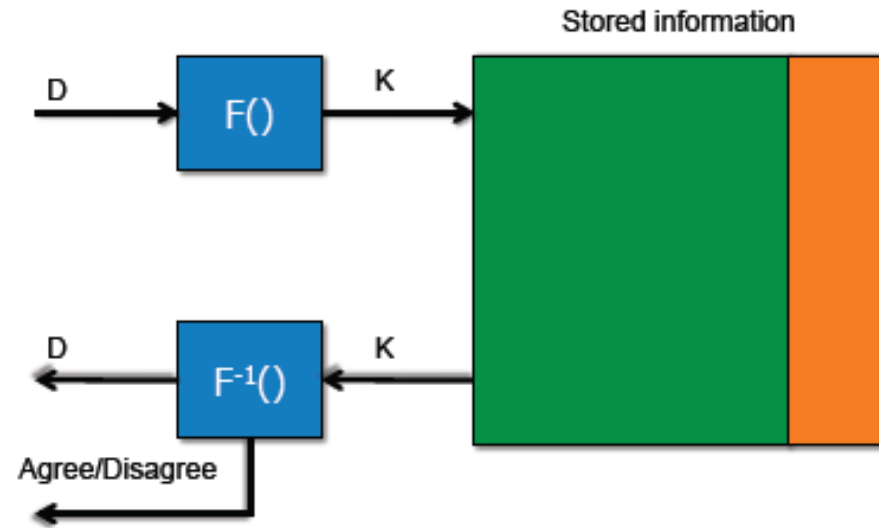
- Security properties:

- Authentication
- Authorization
- Confidentiality
- Integrity
- Availability

# Cryptography

- Encoding operation:  $K=F(D)$ 
  - $D$  is the original data over  $N$  bits
  - $F()$  is the encoding function
  - $K$  is the encoded information, codeword, over  $M>N$  bits
- Decoding operation:  $D=F^{-1}(K)$ 
  - $F^{-1}()$  is the decoding function
- In case there is some error during transmission then codeword  $K$  will be transformed to  $K^*$  and therefore for each  $D$ ,  $K^* \neq F(D)$

# Information redundancy



- Several redundancy schemes are possible that provides:
  - Error detection
  - Error correction
- Schemes differ in the type of error addressed:
  - Single error
  - Multiple error

# Information redundancy

- Examples of information redundancy are:
  - Parity codes
  - Hamming codes
  - Reed Solomon codes

## Parity Codes

- $N$  data bits are coded in a  $N+1$  codeword
  - The codeword is such that:
    - The number of '1' in the codeword is even: even parity
    - The number of '1' in the codeword is odd: odd parity
  - Any single error in the codeword can be detected but not corrected



# References

1. Slides of Dr. Rana Asif Rehman & Dr. Haroon Mahmood
2. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (1994). *Introduction to parallel computing* (Vol. 110). Redwood City, CA: Benjamin/Cummings.
3. Quinn, M. J. *Parallel Programming in C with MPI and OpenMP*, (2003).

## Helpful Links:

1. <https://mpitutorial.com/tutorials/mpi-send-and-receive/>
2. <http://boron.physics.metu.edu.tr/ozdogan/GraduateParallelComputing.old/week11/node2.html> (Odd-Even Sort)
3. <https://mpitutorial.com/tutorials/mpi-scatter-gather-and-allgather/>